

On Component-Based Communication Systems for Clusters of Workstations

Antônio Augusto Fröhlich
Federal University of Santa Catarina
P.O. Box 476
88040-900 Florianópolis - SC, Brazil
guto@inf.ufsc.br

Wolfgang Schröder-Preikschat
University of Magdeburg
Universitätsplatz 2
39106 Magdeburg, Germany
wosch@ivs.cs.uni-magdeburg.de

ABSTRACT

Most of the communication systems used to support high-performance computing in clusters of workstations have been designed focusing on “the best” solution for a certain network architecture. However, a definitive best solution, independently of how well tuned to the underlying hardware it is, cannot exist, for parallel applications communicate in quite different ways. In this paper, we describe a novel design method that supports the construction of run-time systems as an assemblage of components that can be configured to closely match the demands of any given application. We also describe how this method has been deployed in the development of a communication system in the realm of EPOS, a project that aims at delivering automatically generated application-oriented run-time support systems. The communication system in question has been implemented for a cluster of PCs interconnected with Myrinet, and corroborates the effectiveness of the proposed design method.

Categories and Subject Descriptors

D.4.4 [Operating Systems]: Communications Management; D.4.7 [Operating Systems]: Organization and Design—*application-oriented system design*; D.2.13 [Software Engineering]: Reusable Software—*domain engineering*

General Terms

Design, Experimentation, Performance

Keywords

Cluster computing, High-performance computing, Application-oriented operating systems

1. INTRODUCTION

The parallel computing community has been using clusters of commodity workstations as an alternative to expensive parallel machines for several years by now. The results obtained meanwhile, both positive and negative, often lead to the same point: inter-node communication. Consequently, much effort has been dedicated to

improve communication performance in these clusters: from the hardware point of view, high-speed networks and fast buses provide for low-latency and high-bandwidth; while from the software point of view, *user-level communication* [1] enables applications to access the network without operating system intervention, significantly reducing the software overhead on communication. Combined, these advances enabled applications to break the giga-bit-per-second bandwidth barrier.

Nevertheless, good communication performance is hard to obtain when dealing with anything but the test applications supplied by the developers of the communication package. Real applications, not seldom, present disappointing performance figures [12]. We believe the origin of this shortcoming to be in the attempt of delivering generic communication solutions. Most high-performance communication systems are engaged in a “the best” solution for a certain architecture. However, a definitive best solution, independently of how well tuned to the underlying architecture it is, cannot exist, since parallel applications communicate in quite different ways. Aware of this, many communication packages claim to be “minimal basis”, upon which application-oriented abstractions can (have to) be implemented. Once more, there cannot be a best minimal basis for all possible communication strategies. This contradiction between generic and optimal is consequently discussed in [18].

If applications communicate in distinct ways, we have to deliver each one a tailored communication system that satisfies its requirements (and nothing but its requirements). Of course we cannot implement a new communication system for each application, what we can do is to design the communication system in such a way that it becomes possible to tailor it to any given application. In the *Embedded Parallel Operating System* (EPOS) project [5], we developed a novel design method that is able to accomplish this duty. EPOS consists of a collection of components, a component framework, and tools to support the automatic construction of a variety of run-time systems, including complete operating systems.

The particular focus of this paper is on EPOS communication system, which has been implemented for a cluster of PCs interconnected by a Myrinet high-speed network. In the next sections, the *Application-Oriented System Design* method will be introduced, followed by a case study of its applicability to design a communication system. The implementation of this communication system will be discussed later, including a preliminary performance evaluation. The paper is closed with authors’ conclusions.

2. APPLICATION-ORIENTED SYSTEM DESIGN

Application-Oriented System Design (AOSD) is a novel operating system¹ design method that, as the name suggests, is strongly compromised with applications. Its main goal is to produce run-time support systems that can be tailored to fulfill the requirements of particular applications. Accomplishing this task begins with the decomposition of the target domain in abstractions that are natural to application programmers. This is exactly the decomposition strategy promoted by *Object-Oriented Design* [3] and may sound obvious to application designers, but most system designers simply neglect the problem domain and let implementation details, such as target architecture, programming languages, and standardized interfaces, guide the design process [15]. Application programmers, not seldom, get run-time systems that barely resemble the corresponding domain.

The next step is to model software components that properly represent the abstractions from the decomposed domain. Extensive components, that encapsulate all perspectives of an abstraction in a single entity, are not an alternative, since we want components to closely match the requirements of particular applications. A more adequate approach would be to apply the commonality and variability analysis of *Family-Based Design* [14] to yield a family of abstractions, with each member capturing a significant variation and shaping a component. Nevertheless, this approach has the inconvenient of generating a large number of components, thus increasing the complexity of the composition process. We handle this drawback by exporting all members of a family through a single *inflated interface*. In a system designed accordingly, adequate members of each required family could be automatically select by a tool that performs syntactical analysis of the corresponding application's source code.

Another important factor to be considered while modeling abstractions is scenario independence. When a designer realizes, for instance, that a communication mechanism may have to be specialized in order to join a multithreaded scenario, he has to choose between modeling a new family member and capturing this scenario dependency in a separate construct. Allowing abstractions to incorporate scenario dependencies reduces their degree of reusability and produces an explosion of scenario-dependent components. Therefore, an application-oriented design should try to avoid it, only allowing those variations that are inherent to the family to shape new members. The resulting *scenario-independent abstractions* shall be reusable in a larger variety of scenarios, some of them unknown at the time they were modeled.

Scenario specificities, in turn, can be captured in constructs like the *scenario adapters* described in [6]. Because scenario adapters share the semantics of collaborations in *Collaboration-Based Design* [17], one could say that an abstraction collaborates in a scenario. This separation of abstractions and scenario aspects is also pursued by *Aspect-Oriented Programming* [8], nevertheless, though it provides means to support this separation, it does not yet feature a clear domain decomposition strategy.

The primary strategy of Application-Oriented System Design to add functionality to a family of abstractions is the definition of new family members, but sometimes it is desirable to extend the behav-

¹The term “operating system” is used here in its broadest meaning, encompassing all kinds of run-time support systems.

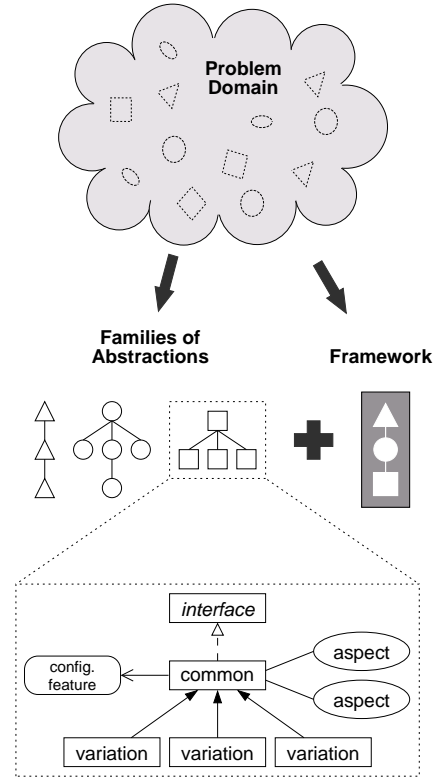


Figure 1: An overview of Application-Oriented System Design.

ior of all members of a family at once. For example, a family of network abstractions might feature multicasting and error detection at application's convenience. Instead of specializing each family member, what would double the cardinality of the family, such extensions could be factorized as *configurable features*. Like scenario aspects, configurable features modify the behavior of all members of a family when activated, but, unlike those, are not transparent. One could say that scenario aspects have “push” semantics, while configurable features have “pull”. A configurable feature encapsulates common data structures and algorithms that are useful to implement a feature in the context of a family, but leave the actual implementation up to each family member. Abstractions are free to reuse, extend, or override what is provided in a configurable feature, but are requested to behave accordingly when the feature is enabled.

After decomposing the problem domain in scenario-independent abstractions and scenario-adapters, organizing the solution domain accordingly becomes straightforward. *Inflated interfaces* hide most details of the solution domain by exporting all members of a family of abstractions, as well as the corresponding scenario adapters, through a single interface. Since these interfaces emanate directly from the problem domain, application programmers should feel comfortable to use them. What is missing to deliver a true application-oriented run-time system is a way to assemble components together correctly and efficiently. By correct assembly we mean preserving the individual semantics of each component in the presence of others and under the constraints of an execution scenario. By efficient assembly we mean preserving their individual efficiency in the resulting composite.

One possibility to produce the desired compositions is to capture a reusable system architecture in a *component framework*. A framework enables system designers to predefine the relationships between abstractions and therefore can prevent misbehaved compositions. Furthermore, a framework defined in terms of scenario adapters (figure 2) can achieve a high degree of adaptability. Efficient composition can be accomplished if the framework uses *Generative Programming* techniques [4], such as *static metaprogramming*. Since static metaprograms are executed at compile-time, a statically metaprogrammed framework can avoid most of the overhead typical of traditional object-oriented frameworks. It is also important to notice that, though component composition would take place at compile-time, nothing would prevent components from using dynamic reconfiguration mechanisms to internally adapt themselves.




Figure 2: A component framework based on scenario adapters.

In brief, *Application-Oriented System Design* is a multiparadigm design method that supports the construction of customizable run-time support systems by decomposing the system domain in families of reusable, scenario-independent abstractions and the corresponding scenario adapters. Reusable system architectures are modeled as component frameworks that can guide the compilation of the target system. Application programmers interact with the system through inflated interfaces, without having to know details about the organization of families or scenarios.

3. THE DESIGN OF AN APPLICATION-ORIENTED COMMUNICATION SYSTEM

We applied Application-Oriented System Design to develop a communication system for clusters of workstations in the realm of project EPOS. By decomposing the domain of high-performance cluster communication, we obtained the families of abstractions shown in figure 3. Application processes communicate with each other using a *Communicator*, which acts as an interface to a communication *Channel* implemented over a *Network*. The messages sent through a *Communicator* can be specified as sequences of bytes of a known length, or they can be covered by an *Envelope*.

3.1 Communicators

A *communicator* is an end-point for a communication channel that enables application processes to exchange data with each other. Therefore, when an application selects a communicator, it implicitly designates the kind of communication channel that will be used. Communicators, like most other EPOS abstractions, are assigned to tasks, thus being shared by their threads. Communicators are real-




Figure 3: Families of abstractions concerning EPOS communication system.




Figure 4: EPOS family of communicators.

ized in EPOS by the *Communicator* family of abstractions shown in figure 4.

The *Link* member of the *Communicator* family realizes an end-point for logical connections between processes that carry byte streams. The *Port* and *Mailbox* members realize end-points for a communication channel in which datagrams flow, but a port always belongs to a single task, while mailboxes can be shared among tasks.

The *ARM_Segment* (*Asynchronous Remote Memory Segment*) member of the *Communicator* family realizes an end-point for a communication mechanism that supports asynchronous access to a memory segment in a remote node. This mechanism is asynchronous because processes manipulating an *ARM_Segment* are not implicitly synchronized and can corrupt the data in that segment. Data read from a remote segment becomes local and private to the reading process. If necessary, synchronization has to be achieved by other means (e.g. distributed semaphores). In order to use this communicator, a process specifies a memory segment on a remote node that has been previously exported by its owner. It can then invoke operations to read from and to write to this segment (asynchronous remote memory segments are not mapped into the address space of processes).

The *AM_Handler* (*Active Message Handler*) member of the *Communicator* family realizes an end-point for active messages [20]. The basic idea behind this concept is that a message, besides transporting data, also carries a reference to a handler that is invoked, in the context of the receiving process, to handle the message upon arrival. This kind of communication is sometimes called single-sided because the receive operation is not explicitly expressible. For this mechanism to work properly, means must be provided to the sending process so it can specify a handler that is valid in the context of the destination process. The most typical answer to this issue is to deploy active messages in an SPMD (*Single Program, Multiple Data*) environment, in which all processes have an equivalent address space. However, indirection mechanisms and the exchange of handler references using other communication mecha-

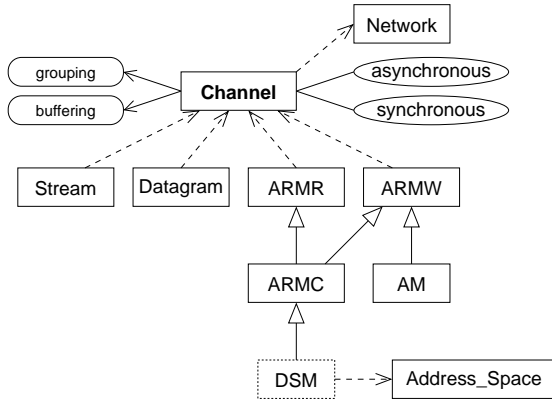


Figure 5: EPOS family of communication channels.

nisms are also possible.

Active messages have been modeled in EPOS in such a way that communication is hidden behind a remote handler invocation, with messages being indirectly exchanged as arguments to the handler. When a process instantiates an `AM_Handler`, it supplies a reference to a handler on a remote process. Afterwards it can invoke the handler supplying arguments that are transparently marshaled in a message and delivered to the remote handler.

The `DSM_Segment` member of the `Communicator` family, which realizes a *Distributed Shared Memory* (DSM) mechanism for EPOS, has been modeled as an extension of the `ARM_Segment` communicator and of a member of the `Segment` family of memory segments. Unlike `ARM_Segments`, however, `DSM_Segments` can be attached to the address space of processes, dispensing with explicit read and write operations. It also comprises mechanisms to grant data coherence. This communicator enables application programmers to write parallel applications for distributed memory machines as if they were shared memory ones.

3.2 Channels

A communication *channel* is the entity effectively responsible for inter-process communication in EPOS. It uses network resources to build a logical communication channel through which messages can be exchanged. A channel implements a communication protocol that, according with the *Basic Reference Model for Open Systems Interconnection* (ISO/OSI-RM) [7], would be classified at level four (transport). EPOS family of communication channels is depicted in figure 5. It was modeled as a dissociated family, whose members are indirectly accessed through the corresponding members of the `Communicator` family.

A communication channel has an implicit capacity. Trying to insert a message into a saturated channel causes the transmitter to wait until the channel can accommodate the message. Likewise, the attempt to extract a message from an empty channel causes the receiver to wait until a message is available. Whether a thread waiting on a channel performs busy or idle waiting hinges on the related configurable feature from the `CPU_Scheduler` family of abstractions. Notwithstanding this, a channel can have its capacity extended by enabling the `buffering` configurable feature. In this case, messages sent through a saturate channel are accumulated for posterior handling.

Sometimes it is desirable to fork a channel, so that a transmitted message is simultaneously multicasted to several receivers, or broadcasted to all receivers. The collective operations used in many parallel applications could be considerably optimized in this way. EPOS allows a channel to be forked when the `grouping` configurable feature is enabled. In this case, special identifiers are supplied to designate a group of communicators as the recipient of a message. The effect of `buffering` and `grouping` configurable features on a channel is illustrated in figure 6.




Figure 6: The effect of configurable features buffering and grouping on communication channels.

The synchronous scenario aspect yields an execution scenario in which the operations used to inject a message into a channel only conclude when the message is extracted from the channel at the receiver's side. In a synchronous communication scenario, processes involved in a message exchange are said to make a "rendezvous". Conversely, the asynchronous scenario aspect modifies these operations so they conclude as soon as the delivery of a message is negotiated. If the `buffering` configurable feature is enabled, this is achieved by copying the message into a buffer and scheduling it for delivery. Otherwise, the sender is supposed not to modify the message until indicated to do so. An operation is provided that enables a process to check for this condition.

The `Stream` member of the `Channel` family realizes a connection-oriented channel that can be used to transfer streams of bytes. It pairs up with the `Link` communicator. The `Datagram` member realizes a channel that supports the transmission of *datagrams*. It has two possible end-points: `Port` and `Mailbox`. Three members concern asynchronous access to a remote memory segment: `ARMR`, `ARMW`, and `ARMC`. They realize communication channels respectively for reading, writing, and copying (reading and writing) from/to a remote memory segment and are delivered to applications through the `ARM_Segment` communicator. The reason to model these three members separately is that read and write operations can be optimized if they do not occur simultaneously.

The `AM` (*Active Message*) channel specializes `ARMW` to introduce the concept of a message handler that is automatically invoked when the message reaches its destination. It pairs up with the `AM_Handler` communicator. The communication channel used to support distributed shared memory would specialize the `ARMC`

channel in order to map it to the address space of processes.

3.3 Networks

A communication channel is, at last, an abstraction of a network, in that networks provide the physical means to build logical channels. The idiosyncrasies of each network technology, however, could require the members of the Channel family to be specialized too often. This picture was prevented in EPOS by modeling networks as members of a uniform family of abstractions, so that all networks are equivalent from the standpoint of channels.

The uniform design of the Network family, which is outlined in figure 7, shall not subdue special features delivered by a particular network, since abstractions in this family implement high-level transport services that are seldom implemented by the hardware. The virtual networks in this family can use special services provided by the network to optimize the implementation of such transport services. Some of these special features are used to implement the configurable features modeled for the family.

The Network family features a member for each network technology supported in the system (e.g. Ethernet and Myrinet). Each member encapsulates a physical network Device. The family also features a Loop device that is used to establish a communication channel between processes executing on the same node. In principle, abstractions in this family are used indirectly through a communicator, but they are also made available for the convenience of applications that need, for instance, to implement special communication protocols.

A set of configurable features, corresponding to operational modes, was modeled for the Network family. These features are interpreted as follows: `ordering` requires messages sent through a network to be delivered at the destination in the same order they were sent; `flow_control` requires a network abstraction to implement flow control; `reliability` requires a network to assure error-free delivery of messages; `broadcast` enables the interpretation of broadcast addresses, so messages can be broadcasted to all hosts in the local network; `multicast` enables the interpretation of multicast addresses, causing a message to be delivered at multiple hosts. These configurable features are usually specialized for each family member to profit from eventual hardware support.

3.4 Message Envelopes

The members of the Communicator family can be used to exchange unstructured messages in the form of sequences of bytes of a certain length. However, it might be adequate for some applications to count on an *envelope* abstraction to cover a message before sending it. Such an envelope would be allocated from the operating system, loaded with one or more messages, and then inserted into a communication channel through a communicator. An




Figure 7: EPOS family of networks.

envelope allocated by the operating system would enable several optimizations, ranging from cache alignment to zero-copy processing. Besides, additional information can be put in the envelope to describe and protect messages. After all, an envelope would enable a comfortable syntax to express communication in object-oriented applications, for example:

```
Envelope envelope(recipient, length);
envelope << "Hello world!";
communicator << envelope;
```

EPOS supports the concept of *message envelope* through the Envelope uniform family of abstractions represented in figure 8. The maximum length of message that an envelope can hold is specified when it is instantiated, while the effective length of the message(s) it contains is dynamically determined. An envelope must be addressed before it is posted.

The Envelope family comprises two members: `Untyped` and `Typed`. The former realizes a simple message envelope that can be used to gather messages before sending, while the latter collects type information for each message inserted to enable format conversions on heterogeneous systems. A *secure envelope* was not modeled due to the characteristics of a dedicated computing system, which usually do not require encryption nor authentication of messages.

3.5 Scenario Aspects

Application-Oriented System Design is particularly concerned with scenario independence. When a domain entity is identified, considerations about its origin are made in order to decide whether it will shape an abstraction or a scenario aspect. EPOS scenario aspects were modeled in accordance with this principle, yielding reusable pieces of software that can be controllably applied to the system abstractions described in the previous section. Figure 9 shows scenario aspects that apply to EPOS communication system.

Identification: instances of EPOS abstractions can be identified in four ways: (a) by its address in memory in single-task scenarios; (b) by a pair (`type`, `unit`) in stand-alone multi-tasking scenarios; (c) by a tuple (`host`, `type`, `unit`) in distributed scenarios; and (d) by a sparse capability of the type (`host`, `type`, `unit`, `rand`) for protected distributed scenarios. Although embedding location in identifiers complicates the migration of objects in a distributed environment, EPOS opted for this solution because it is efficient and fulfills the demands of most parallel systems.

Sharing: EPOS abstractions can be shared both in local and distributed environments. The family of sharing scenario aspects features two members: the first performs reference




Figure 8: EPOS family of message envelopes.

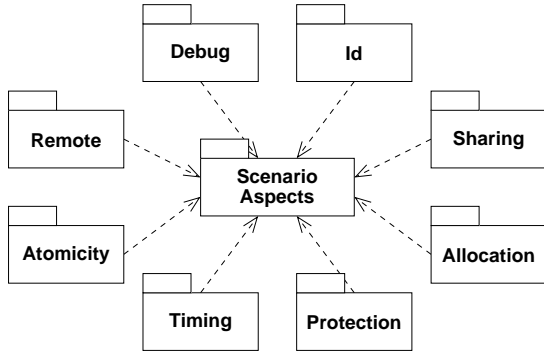


Figure 9: Families of scenario aspects concerning EPOS communication system.

counting, while the second extends the family to support possession lists. When one of these aspects is enabled, system abstractions gain two extra constructors that are used to designate sharing. The first is the ordinary copy constructor, which takes a reference to an existing system object as argument and returns a share to that object. This constructor is restricted to share abstractions inside a single address space. The second constructor takes an identifier as argument and therefore can be used independently of locality, inclusive to share remote objects. A C++ program could deploy these constructors as follows:

```
Abstraction instance;
Abstraction share1(instance);
Abstraction share2(share1.id());
```

Allocation: in order to reduce run-time overhead, EPOS abstractions can be allocated in advance during system initialization, so that applications get all resources they will need by the time they start execution. This is accomplished by a scenario aspect that interacts with EPOS initialization programs to pre-allocate vectors of objects that are controlled by special allocators. Otherwise, a counterpart scenario aspect takes care for dynamic allocation.

Protection: EPOS supports a protected execution scenario for applications. This scenario profits from the capability identification aspect to build an access control mechanism similar to the one described in [11].

Timing: two scenario aspects concerning timing have been modeled for EPOS communication system. One supports the specification of time-outs for abstractions' operations; the other inserts delays on operations to balance execution time in heterogeneous systems.

Atomicity: when multiple threads are allowed to execute in parallel, it is possible that they simultaneously "enter" the operating system, i.e. pass to execute operating system code to accomplish a system service. EPOS handles the synchronization pitfalls brought about by reentrance ensuring that system operations are *atomic*. In this way, transformations of the state of system objects either occur completely or do not occur at all.

Remote Invocation: in a distributed environment, processes may need to access system resources residing on remote nodes.

In order to do so, an application would have to create a process on each node containing useful resources and deploy a *Communicator* to interact with them. However, EPOS eliminates the burden of explicit message passing for accessing remote resources with a *Remote Object Invocation (ROI)* mechanism similar to the one described in [9]. Such a mechanism hides inter-process communication behind ordinary method invocations, so that processes can transparently access remote resources. This scenario aspect that can be transparently applied to virtually any abstraction.

Debugging: being able to trace the invocation of system operations, or to watch the state of system abstractions, can be useful to debug application programs. Likewise, being able to summarize how much time an application spends with each system abstraction may be a source of optimization. EPOS supports these features through a family of scenario aspects.

3.6 System Generation

When completely implemented for a variety of network architectures, EPOS communication system will yield a large number of components that will be stored in a repository together with several other subsystems. With such a large number of components, selecting and configuring the right ones in order to produce an application-oriented system may become a defying activity, even when assisted by visual tools. Hence, Application-Oriented System Design proposes all members of a family to be exported through a single, inflated interface. In this way, application programmers can design and implement their applications referring to fewer interfaces and ignoring the particular properties of each component. Actually, the programmer catches a comprehensive perspective of the family, as though a super-component was available, and uses the operations that better fulfill his requirements². As an example, the inflated interface of the *Communicator* family is depicted in figure 10.

The process of binding an inflated interface to one of its realizations can be automated if we are able to clearly distinguish one realization from another. In EPOS, we identify realizations through the signatures of their methods, so that syntactical analysis of application source code can identify which of the realizations are needed. If two realizations present the same set of signatures, as with *Port* and *Mailbox* in figure 10, syntactical analysis might not be enough to decide for one of them, and user intervention may be required. Nevertheless, although *Port* and *Mailbox* differ only semantically³, the syntactical analysis of other components may render one possibility invalid. For example, if the application is known to execute on a single-task-per-node basis, a scenario with multiple receivers is not possible, breaking the tie in favor of *Port*.

The set of selected family members, in addition to information obtained from the user, defines the execution scenario for the application. As proposed by Application-Oriented System Design, scenario peculiarities are applied to abstractions by means of scenario adapters. In EPOS, a scenario adapter wraps an abstraction as to enclose invocations of its operations between the *enter* and *leave*

²In case an application programmer with enough expertise about the system wishes to extend a component, or bypass automatic configuration, the individual interfaces of each family member are also made available.

³Both *Port* and *Mailbox* support multiple senders, but the first supports a single receiver, while the second support multiple receivers too.

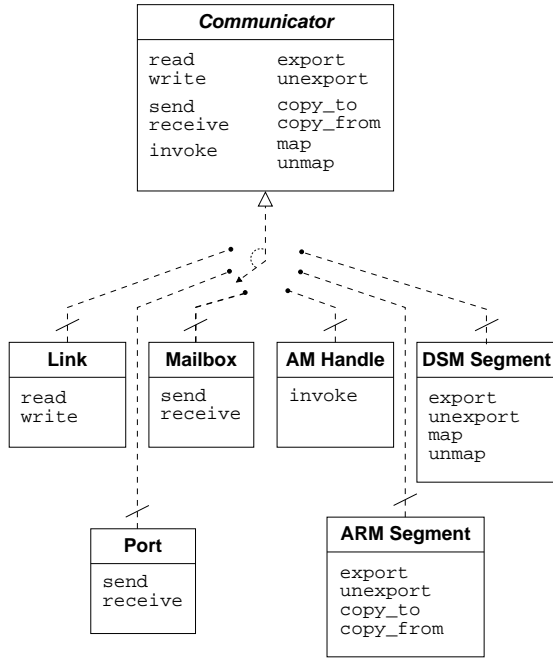


Figure 10: The Communicator inflated interface and its realizations.

scenario primitives (see figure 11). Besides enforcing scenario specific semantics, a scenario adapter can extend the state and behavior of an abstraction, for it inherits from both scenario and abstraction. For example, all abstractions in a scenario could be tagged with a capability, without internal modifications, by associating the capability with the corresponding scenario.

EPOS statically metaprogrammed framework is defined around a collection of interrelated scenario adapters. As shown in figure 11, scenario adapters are designed as parametrized classes that take a component (abstraction) as parameter. Hence they can act as placeholder for components in the framework. In order to generate a system, information about the mapping of inflated interfaces to realizations, and also about system-wide properties such as target ar-

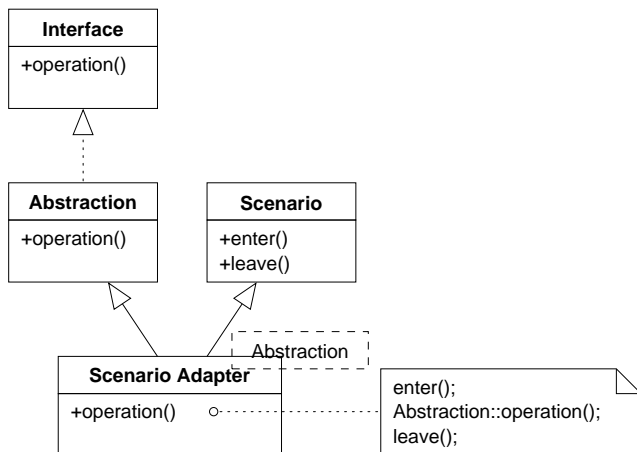


Figure 11: The structure of a scenario adapter.

chitecture and protection, are passed as input to the metaprogram. The resulting system would include only the components needed to support the corresponding application in the respective execution scenario.

4. THE IMPLEMENTATION OF THE COMMUNICATION SYSTEM

The design of EPOS described above is being implemented as a collection of components, a framework, and a set of tools that support automatic generation of application-oriented run-time systems. Currently the system can run in two modes: native on IX86 computers, and guest on LINUX systems. The IX86-native version can be configured either to be embedded in the application (single-task), or as μ -kernel (multi-task), while the LINUX-guest version comprises a library and kernel loadable modules.

Components are being implemented in C++ and described in XML. The XML description is used by the tools that support automatic system generation. The framework is also implemented in C++, but mainly with its built-in static metalanguage. The tools to proceed syntactical analysis of applications, to configure the target system, and to check configuration dependencies are made available to users through a compiler wrapper similar to `mpicc`. This enables users to implicitly generate the run-time system during the compilation of applications. Nevertheless, if these tools fail to configure the system, user intervention is requested via an interactive graphical tool that supports configuration adjustments by feature selection⁴.

EPOS family of communication abstractions is currently being implemented for the Myrinet high-speed network [2]. So far, we concluded the implementation of the Port and Mailbox Communicators, the Datagram Channel, the Myrinet Network, and a mechanism to support *remote object invocation* (ROI). These components can be adapted to the following scenarios: Protected, Multitask, Multithread, and Global. The Protected scenario ensures that only authorized agents gain access to abstractions. The Multitask and Multithread scenarios adapt abstractions to execute in the presence of multiple tasks and threads. The Global scenario adapts abstractions to interact in a cluster-wide environment of active objects, hiding communication behind ordinary method invocations.

With these components we generated a couple of application-oriented run-time systems. One of them supports two simple applications that communicate intensively in a producer/consumer fashion. For this purpose, they use the Port Communicator, the Datagram Channel, and the Myrinet Network. Figures 12 and 13 show respectively the latency and the bandwidth available to these applications in both IX86-native and LINUX-guest modes. The hardware test-bed for this measurements consisted of two PCs connected to the same Myrinet switch. Each PC has a 266 MHz Pentium II processor, 128 MBytes of memory (10 ns DRAM) on a 66 MHz bus, and a 32-bits Myrinet NIC on a 33 MHz PCI bus.

The difference in favor of the IX86-native version arises from the contiguous memory allocation method adopted, which allows the DMA engines on the Myrinet card to be programmed with logical addresses and eliminates an additional message copy into a system DMA buffer. This difference could have been even more expressive if the applications were multithreaded, since the extra copy

⁴The same tool can be used to tailor the system manually.

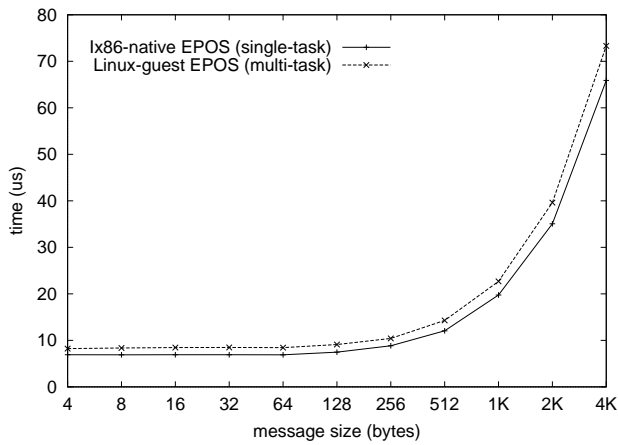


Figure 12: Port/Datagram/Myrinet one-way latency.




Figure 13: Port/Datagram/Myrinet one-way bandwidth.

would have concurred with application threads for processor time and especially for memory bandwidth. Nevertheless, most parallel applications execute on a single-task-per-node basis and will benefit from the single-task versions of EPOS. Other communication systems, such as Active Messages [10], Fast Messages [13], PM [19], and BIP [16], run exclusively on top of ordinary operating systems, such as UNIX or WINDOWS NT, and have no alternative to escape the extra copy than making a system call to translate logical addresses into physical ones, what is usually even more time consuming⁵.

Furthermore, EPOS quality evaluation should not be restricted to performance. Because only the components effectively required by the application are included, the resulting system is usually extremely compact. The system in the example above, which in addition to communication also includes process and memory management, has a size of 11 KBytes. This means less resource consumption and less space for bugs. Usability is also improved, since EPOS visible interfaces are defined in the context of applications.

⁵Sharing system DMA buffers with applications is not really an alternative, because they are usually restricted in size and will not be able accommodate the large data structures typical of parallel applications. It would most likely result in the application performing the additional copy.

5. CONCLUSIONS

In this paper we introduced *Application-Oriented System Design*, a novel design method that prevents the monolithic conception of generic solutions that fail to scale along with application demands. We also described how this method has been deployed to construct a communication system for the Myrinet high-speed network. This communication system, implemented in the realm of project EPOS, consists of a collection of *application-ready, scenario-independent abstractions* (components) that can be adapted to specific execution scenarios by means of *scenario-adapters* and can be arranged in a *statically metaprogrammed framework* to produce an application-oriented communication system. The system is presented to application programmers through *inflated interfaces* that gather all variations of an abstraction (family members) under a single comprehensive interface. By programming based on these interfaces, programmers enable EPOS tools to automatically generate an adequate system for their applications.

The results obtained so far are highly positive and help to corroborate the guidelines of *Application-Oriented System Design*, as well as EPOS design decisions. The evaluation of EPOS communication system revealed performance figures that, as far as we are concerned, have no precedents in the history of PC clusters interconnected with 32-bits Myrinet. Nevertheless, EPOS is a long term project that aims at delivering application-oriented run-time systems to a large universe of applications. Therefore, several system abstractions, scenario adapters, and tools are still to be implemented or improved.

6. ACKNOWLEDGMENTS

The research work described in this paper was carried out at the Fraunhofer FIRST Institute in the realm of the SNOW project. Additional funding has been received from the CAPES Foundation (grant no. BEX 1083/96-1) and the Deutsche Forschungsgemeinschaft (grant no. SCHR 603/1-1).

7. REFERENCES

- [1] R. A. F. Bhoedjang, T. Rühl, and H. E. Bal. User-level Network Interface Protocols. *IEEE Computer*, 31(11):53–60, Nov. 1998.
- [2] N. J. Boden, D. Cohen, R. E. Felderman, A. E. Kulawik, C. L. Seitz, J. N. Seizovic, and W.-K. Su. Myrinet: A Gigabit-per-Second Local-Area Network. *IEEE Micro*, 15(1):29–36, Feb. 1995.
- [3] G. Booch. *Object-Oriented Analysis and Design with Applications*. Addison-Wesley, 2 edition, 1994.
- [4] K. Czarnecki and U. Eisenecker. *Generative Programming: Methods, Tools, and Applications*. Addison-Wesley, 2000.
- [5] A. A. Fröhlich and W. Schröder-Preikschat. Tailor-made Operating Systems for Embedded Parallel Applications. In *Proceedings of the 4th IPPS/SPDP Workshop on Embedded HPC Systems and Applications*, volume 1586 of *LNCS*, pages 1361–1373, San Juan, Puerto Rico, Apr. 1999. Springer.
- [6] A. A. Fröhlich and W. Schröder-Preikschat. Scenario Adapters: Efficiently Adapting Components. In *Proceedings of the 4th World Multiconference on Systemics, Cybernetics and Informatics*, Orlando, U.S.A., July 2000.
- [7] International Organization for Standardization. *Open Systems Interconnection - Basic Reference Model*, Aug. 1981. ISO/TC 97/SC 16 N 719.

- [8] G. Kiczales, J. Lamping, A. Mendhekar, C. Maeda, C. V. Lopes, J.-M. Loingtier, and J. Irwin. Aspect-Oriented Programming. In *Proceedings of the European Conference on Object-oriented Programming '97*, volume 1241 of *LNCS*, pages 220–242, Jyväskylä, Finland, June 1997. Springer.
- [9] H. M. Levy and E. D. Tempero. Modules, Objects, and Distributed Programming: Issues in RPC and Remote Object Invocation. *Software — Practice and Experience*, 21(1):77–90, Jan. 1991.
- [10] S. S. Lumetta, A. M. Mainwaring, and D. E. Culler. Multi-Protocol Active Messages on a Cluster of SMP's. In *Proceedings of Supercomputing '97*, Sao Jose, USA, Nov. 1997.
- [11] S. J. Mullender and A. S. Tanenbaum. The Design of a Capability-based Distributed Operating System. *The Computer Journal*, 29(4):289–300, 1986.
- [12] Numerical Aerospace Simulation Systems Division. *The NAS Parallel Benchmarks*, online edition, Nov. 1997. [<http://www.nas.nasa.gov/Software/NPB/NPB2Results/index.html>].
- [13] S. Pakin, V. Karamcheti, and A. A. Chien. Fast Messages: Efficient, Portable Communication for Workstation Clusters and Massively-Parallel Processors. *IEEE Concurrency*, 5(2), June 1997.
- [14] D. L. Parnas. On the Design and Development of Program Families. *IEEE Transactions on Software Engineering*, SE-2(1):1–9, Mar. 1976.
- [15] R. Pike. Systems Software Research is Irrelevant. Online, Feb. 2000. [<http://cm.bell-labs.com/who/rob/utah2000.ps>].
- [16] L. Prylli and B. Tourancheau. BIP: a New Protocol Designed for High Performance Networking on Myrinet. In *Proceedings of the International Workshop on Personal Computer based Networks of Workstations*, Orlando, USA, Apr. 1998.
- [17] T. Reenskaug, E. P. Andersen, A. J. Berre, A. Hurlen, A. Landmark, O. A. Lehne, E. Nordhagen, E. Næss-Ulseth, G. Oftedal, A. L. Skaar, and P. Stenslet. OORASS: Seamless Support for the Creation and Maintenance of Object-oriented Systems. *Journal of Object-oriented Programming*, 5(6):27–41, Oct. 1992.
- [18] W. Schröder-Preikschat. *The Logical Design of Parallel Operating Systems*. Prentice-Hall, Englewood Cliffs, U.S.A., 1994.
- [19] H. Tezuka, A. Hori, Y. Ishikawa, and M. Sato. PM: An Operating System Coordinated High Performance Communication Library. In *High-Performance Computing and Networking*, volume 1225 of *LNCS*, pages 708–717, Apr. 1997. Springer.
- [20] T. von Eicken, D. E. Culler, S. C. Goldstein, and K. E. Schauer. Active Messages: a Mechanism for Integrated Communication and Computation. In *Proceedings of the 19th Annual International Symposium on Computer Architecture*, pages 256–266, Gold Coast, Australia, May 1992.