

Blue Flash Memory Technology

Antônio Augusto Fröhlich
Marcelo Trierveiler Pereira
Arliones Stevert Hoeller Junior
Felipe Zimmermann Homma

6th May 2003

Chapter 1

Flash Memory Technology

This chapter brings an overview of Flash Memory Technologies, aiming at establishing the state-of-the-art in the field and serving as a basis for the forthcoming sections of this text.

Flash memory is a solid-state, non-volatile, re-writable memory that works like a *Random Access Memory* [RAM] unit and a hard disk drive combined. Flash memory stores bits of electronic data in memory cells, just like DRAM, but it also works like a hard-disk drive in that when the power is turned off, the data remains in memory. Because of its high speed, durability, and low voltage requirements, flash memory is ideal for use in many applications, such as digital cameras, cell phones, printers, handheld computers, pagers, and audio recorders [10].

1.1 General Concepts

Flash Memories are very similar to *Electrically Erasable Programmable Read-Only Memories* [EEPROM], the main difference being that flash memories can only be erased in chunks (called sectors), and this is the origin of its name (erase sectors in a “flash”). Such erasing scheme simplifies the circuitry, allowing for a greater density in regard to an equivalent EEPROM. Actually, flash memories can achieve densities similar to EPROMs.

There are several technologies a flash memory cell can be built on: NOR, DINOR, T-Poly, AND, NAND; each of which requires a particular programming and erasing method [37]. These technologies allow flash memories to retain stored data without a permanent power supply for periods as long as twenty years. Nevertheless, the very same technologies are responsible for one of the biggest shortcomings of flash memories: the limitation in the number of write/erase operations that can be performed before material fatigue becomes critical and compromises data consistency. Notwithstanding, the typical rewrite limit (in the order of hundred thousands) is acceptable for many applications.

Since flash memories can only be erased in sectors, the size of a sector becomes a crucial operational factor for such memories. In order to allow for a more rational use of the memory available in a flash, some models provide sectors of different sizes and properties. For instance, a smaller “boot sector” can be provided that supports write-protection, or yet an special sector can hold a unique serial number assigned at manufacturing-time.

1.2 Architectural Overview

Flash, EPROM, and EEPROM devices use basically the same floating gate mechanism to store data, but they deploy distinct reading and writing mechanisms [17]. In all cases, the basic memory cell consists of a single *MOS transistor* [MOSFET] with two gates: a *control gate* that is connected to the read/write control circuitry and a *floating gate* that is located between the control gate and the channel — the

part of the MOSFET through which an electric current can flow between the so-called *source* and *drain* terminals.

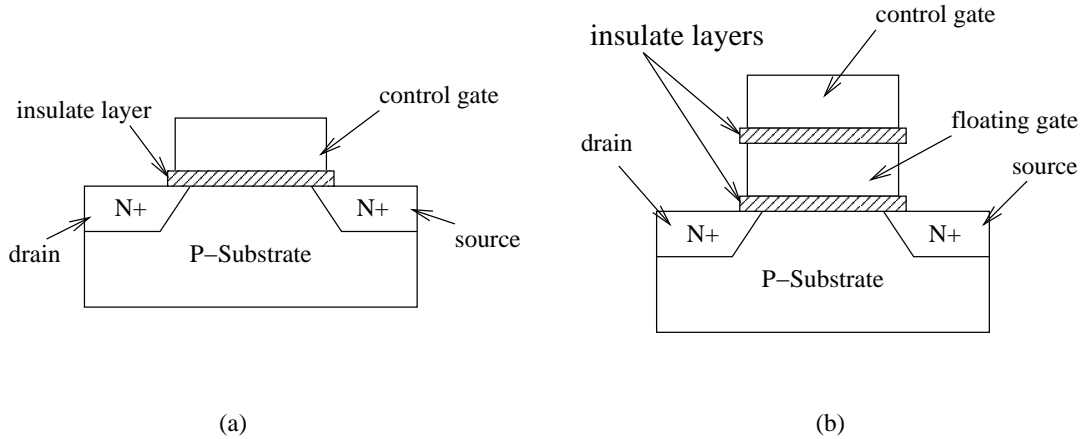


Figure 1.1: schematic of (a) MOSFET and (b) floating-gate transistor.

In a standard MOSFET, a single *gate* terminal controls the electrical resistance of the channel: an electric voltage applied to the *gate* controls how much current flows between *source* and *drain*. Figure 1.2 illustrates the behavior of such a kind of transistor. Current is only present in the channel if a voltage is applied to the control gate. The voltage needed to build up the “bridge” over the channel is commonly called *Threshold Voltage* (V_t). The MOSFETs used in non-volatile memories include a *second gate* that is completely surrounded by an insulating layer, that is, it is electrically insulated from the rest of the circuitry (see figure 1.1).

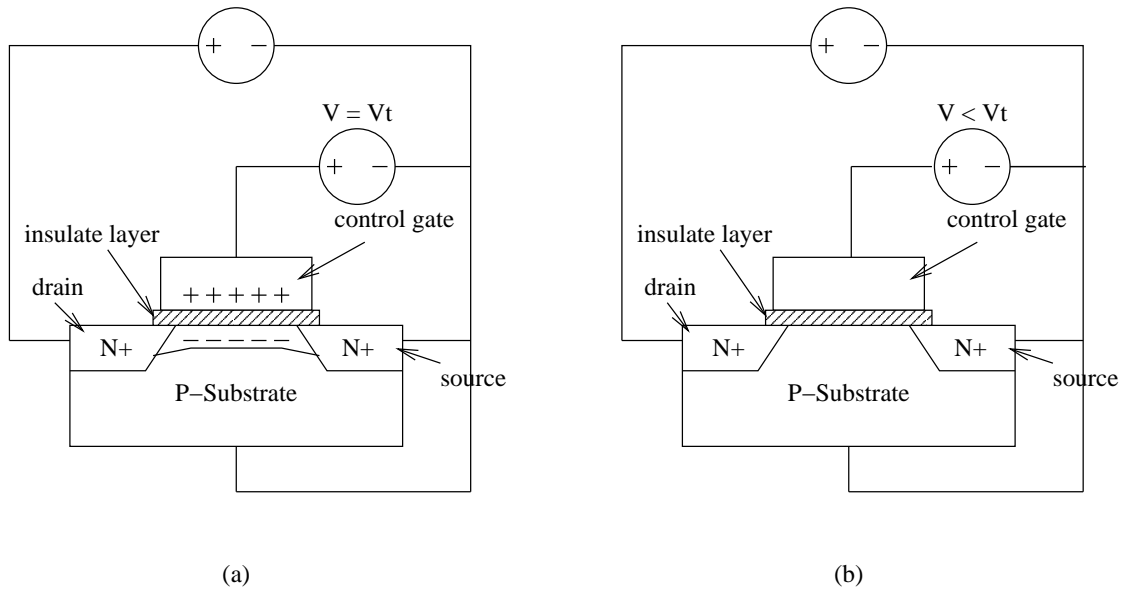


Figure 1.2: MOSFET behavior: with (a) and without (b) current.

Because the floating gate is physically very close to the MOSFET channel, even a small electric charge has an easily detectable effect on the electrical behavior of the transistor. By applying appropriate signals to the control gate and measuring the change in the transistor behavior, it is therefore possible to determine whether or not there is an electric charge on the floating gate (see figure 1.3). In a floating-gate transistor, V_t will be linked with the number of electrons trapped in the floating-gate. It will be as high as the number of electrons trapped in the floating-gate. Similarly to a MOSFET, the flow in the channel is from source to drain.

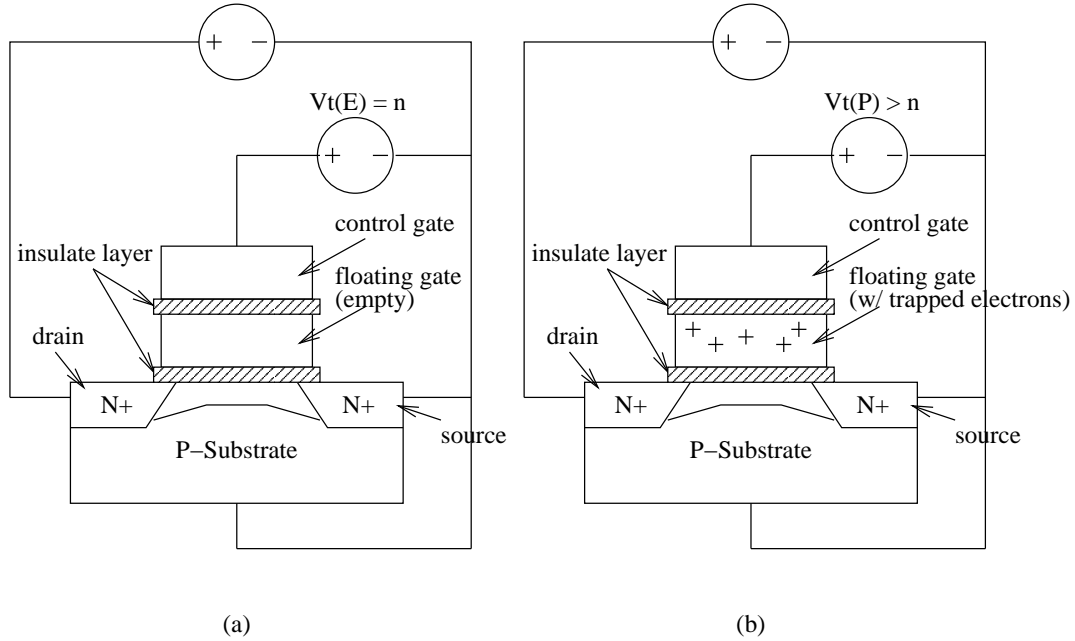


Figure 1.3: Floating-gate transistor behavior: (a) without and (b) with trapped electrons.

Since the floating gate is electrically insulated from the rest of the transistor, special techniques are required to move electrons to and from it. One technique consists in filling the MOSFET channel with high-energy electrons by applying relatively high voltages to the control gate and the drain of the MOSFET. Some of these "hot" electrons have sufficient energy to cross the barrier between the channel and the floating gate. When the high voltages are removed, these electrons remain trapped on the floating gate. This is the method used to program a memory cell in EPROM and flash memories.

This technique, known as *Channel Hot Electron* [CHE] injection, can be used to load an electric charge onto the floating gate, but does not provide a way to discharge it. In order to discharge a floating gate, flash memories¹ tackle on a quantum effect known as *tunneling*: electrons are removed from the floating gate by applying a voltage to the MOSFET that is large enough to cause electrons to 'tunnel' across the insulating layer [30].

Traditionally, the floating gate mechanism has been used to store a single data bit, which is read by comparing the MOSFET threshold voltage with a reference value. More sophisticated techniques make it possible to distinguish more than two floating gate charge states, thus enabling two or more bits to be stored on a single floating gate. This is an important technology breakthrough, because storing two bits/cell doubles the memory capacity for a given cell size [14].

1.3 Operation

There are three basic operations that can be performed on a flash: *read*, *program*, and *erase*. Reading and programming a flash is usually as fast as the equivalent operations on a DRAM (read and write), but the erasing operation is much slower than reading or programming. This shortcoming is being tackled through partitioning and also through the inclusion of pause operations that allow for operating interweaving in a single partition.

The execution of these operations is managed and validated by the flash own state-machine. Usually, the state-machine of a flash is also able to detect operation time-outs, which may be an indicative of cell

¹In EPROM memories, the floating gate is discharged by flooding the entire memory array with ultra-violet light — the high-energy light penetrates the chip structure and impart enough energy to the trapped electrons, allowing them to escape the floating gate.

degeneration.

Reading: reading from a flash memory is similar to reading from traditional memory devices (e.g. DRAM). In order to reduce latency and improve bandwidth, some flash memory units deploy sophisticated access modes: besides traditional asynchronous bit and word read modes, they often support asynchronous page read (buffering a whole page to speedup subsequent accesses to the same page) and synchronous burst read (multiple data words from a single address).

Figure 1.4 illustrates the procedure of reading a flash cell. Applying a reading voltage ($V_t(R)$) to the control gate can produce one of two situations: (a) if the floating-gate is empty, current will be detected in the channel; or (b) if there are electrons in the floating-gate, then no current will be formed. It is also possible to define several levels for V_t by controlling the number of electrons in the floating-gate. This is the principle of multilevel cells, which rely on different states to store more than a bit per transistor.

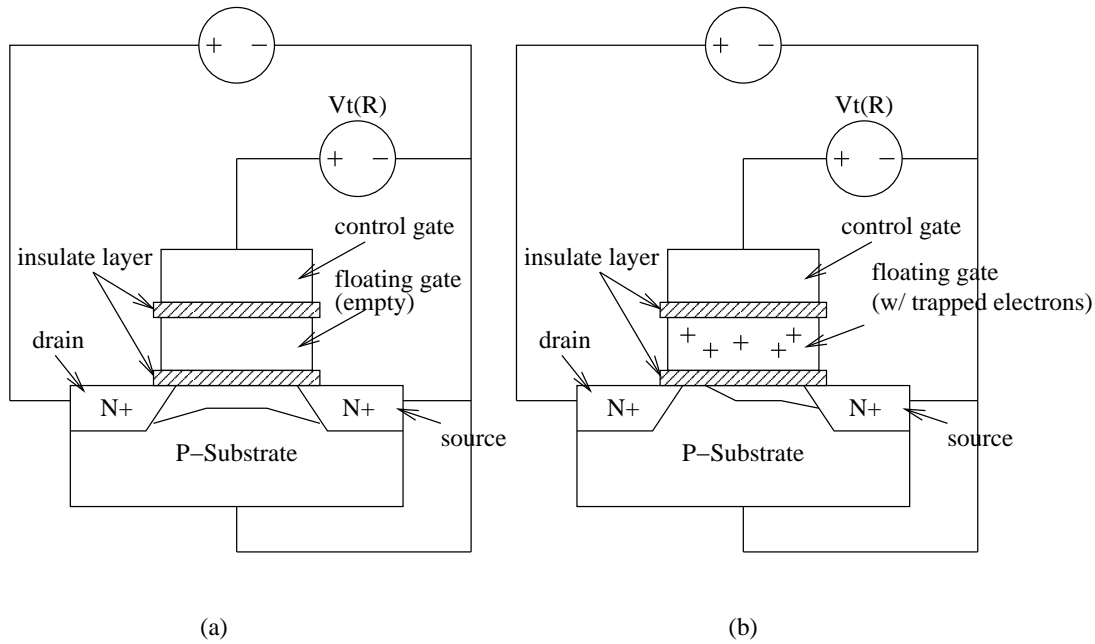


Figure 1.4: Reading operation on a floating-gate transistor.

Programming: a typical flash memory is delivered with all sectors erased, i.e. with all bits set to “1”. Thus, programming (writing to) a flash consists in turning some “1s” into “0s”. Note that the reverse operation, turning “0s” into “1s”, is not defined². The only way to set a bit in a flash is by erasing the whole sector. Some flash models support programming single bits, words, and even blocks (burst write). Some models also provide buffers to speed up programming, thus supporting a write operating similar to DRAM (as long as one does not try to overwrite a “0” with an “1”). At transistor level, programming a flash consists in injecting electrons into the floating gate.

Erasing: Erasing a sector of a flash memory unit, i.e. setting all of its bits to “1”, is achieved by removing electrons from floating gates. Depending on the technology used, a flash memory may be subjected to the so called “over-erasing” phenomenon: erasing a cell whose value is already “1” puts that cell in a state that prevents further programming. Flashes exposed to the phenomenon usually handle it internally by programming all bits of a sector (setting to “0”) before erasing them (setting to “1”). Therefore, erasing a flash becomes a trivial operation for firmware/driver programmers.

²Trying to program a previously zeroed bit of a flash to “1” usually has no effect, though it probably causes the flash’s state-machine to signalize a condition flag that might trigger external events.

1.4 On the Market

Since the invention of the flash memory, manufacturers have been looking for alternatives to increase performance and capacity. As flash memories gain new markets, technologies are being incorporated to turn it a more reliable choice over other persistent memory technologies. Some of the most significant improvements now available on the market are:

Non-Uniform Sector Size Architecture: some models are designed with non-uniform sector size. They allow for specialized sectors such as boot sectors and sensitive data without waste of space.

Common Flash Interface [CFI]: specifies a standard command set for flash memories. It has been pushed by market leaders in an attempt to make the development of flash software vendor independent.

Multi-partitions Architecture: multi-partition or multi-bank architecture supports multiple operations to be simultaneously performed on the same flash memory unit, thus allowing for parallel operation.

Multilevel Cell Technology: multilevel cell intends to raise the density (i.e. amount of bits) of the flash without raising the die area. This is reached by controlling the amount of electrons that are trapped in the floating-gate as discussed earlier in this section. Typical models of multilevel cell flash memories store 2 bits per cell.

Security: some models include security registers —one-time programming sectors or sectors dedicated to store security data (like unique identifier numbers).

1.4.1 Case Studies

The forthcoming discussion presents several flash memory technologies currently available on the market. The purpose of this section is to identify each technology highlights.

AMD

<http://www.amd.com/>

AMD delivers various CFI compliant models, with operational voltages ranging from 1.8V to 5.0V, densities up to 256 Mb and with boot sectors (top or bottom configurations). Their products operate in temperatures from commercial (0 — +70° C) to super-extended (-55 — +145° C).

They have models with multi-partitions, MirrorBit technology and Dual Operation (Flash + SRAM hybrid). For safety reasons, their products are able to lock one dedicated sector.

Dual Operation: by incorporating a SRAM unit AMD delivers a “hibrid” flash model that can make multiple operations without partitioning the memory. All operations are executed on the SRAM unit and later committed to flash.

MirrorBit Technology: AMD multi-bit technology. This technology differs from the other manufacturers by implementing two distinct bits instead of 2^2 states to create two bits. This way they can offer a high density model without compromising performance or reliability. This is done by building a symmetrical transistor with the sources and drain replicated in both sides (see Figure 1.5), and with a very careful and controlled injection of electrons in one side of the floating gate.

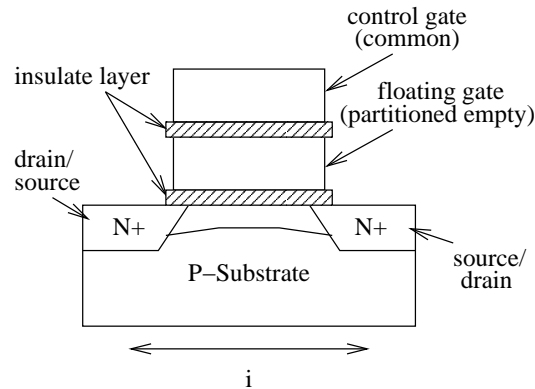


Figure 1.5: MirrorBit transistor schematic.

Intel

<http://www.intel.com/>

Intel products operates in voltages from 1.8 V to 5 V, some models have higher output voltages for compatibility purposes; they have multi-partition, non-uniform size sectors (boot sectors in top or bottom configuration) and multilevel cell technology (StrataFlash models), all CFI compliant.

Security mechanism is a bunch of bits composed by two parts: one programmed by Intel and other programmed by the user.

Besides the BCS, Intel have the Extended Command Set (ECS) CFI compliant, which makes use of their own hardware character.

Enhanced Factory Programming (EFP) and Buffered Enhanced Factory Programming (BEFP): these two Intel exclusive technologies are used to minimize the flash programming in controlled environments, usually to pre-program the chip and put it in the equipment later. This is made by turning all data checks off during the program stage, verifying all the correctness later. Both technologies are the same, but one applies to standard flash memories and other to Strata (multilevel) flash memories.

ST Semiconductors

<http://us.st.com/>

ST have multi-level cell technology, multi-partitions, boot sector (top or bottom configurations), programmable security sector and CFI compliance. Voltages from 1.8 to 5.0V, temperatures from commercial to super-extended.

ST Semiconductors has a range of specialized products to the automotive market. Their M29F (5V) and M29W (3V) families can operate in automotive environment (-40° to +125° C); M58BF008 (8Mb)

and M58BW016 (16Mb) models are dedicated to this application class.

ATMEL

<http://www.atmel.com/>

ATMEL flash memories operate in voltages from 2.7 to 5V, 512 to 32 Mbits densities, concurrent Read & Write, boot sector, multi-partitions and frequencies reaching 100 MHz. In some models it is possible to define 16 or 32 bits databus.

Fast Programming Time: if there is high voltage available (12V), it is possible to reach very high programming speeds (30us per word/double word).

Serial Flash: serial flash memory uses Serial Peripheral Interface (SPI), and it can be used as replacement for SPI compliant EEPROMs, without changes of the board layout (if pinage also compatible). It operates at 20 MHz and in low voltages (2.7—3.6 V), densities ranging from 512 Kbits to over 4 Mbits.

Data Flash: SPI compliant hybrid serial/parallel flash memory.

MICRON

<http://www.micron.com/>

Micron products are CFI compliant and have their own extensions (SCS). Their products have dedicated boot sector (in top or bottom configuration), multi-partition architecture, operational voltages from 2.7 to 5.0V, select databus size (8 or 16 bits), temperatures varying from commercial to extended.

SynchFlash: this family, a SDRAM similar interface (that can imitate a SDRAM chip) is introduced. It has high read performance (equals SDRAM read performance) consequently turned the chip into a very competitive choice for execute-in-place applications. This Family is second sourced by ATMEL. There are plans (in a near future) to reach the same read performance of DDR SDRAM.

1.5 Future of Non-volatile Memory Technologies

Today flash memory is the best choice for fast and small persistent storage due to its high speed and density, but new technologies are coming, specially those utilizing magnetic media. There are also other researches in optics and nanotechnology areas.

The new material and new research tools, make the creation of new storage medias and devices possible. The near future promises faster and more dense devices than flash memory.

By now ferroelectric and magneto-resistive memory have great potential to be deployed with great advantage over flash memory, because these technologies are not based on read-only schema. For future of holographic devices, we can expect a massive storage at great speeds. IBM has been studying nanotechnology since 1990 and the prototypes are promissing.

1.5.1 Future of Flash

Flash memory will be faster and more dense and reliable. Read speed increases are visible even these days, by buffering, paging and other indirect access techniques. But flash devices will be limited by integrated circuit technology and material limits.

Advanced manufacturing techniques can reduce the size of the tracks and transistors giving flash memory more density, raising the write speed, but its scale limits are bound to material (silicon) limits.

Unfortunately this technology was based on read-only memory technology and thus, writing will continue to be the bottleneck, and maybe the hybrid systems (embedded RAM + Flash) will partially manage this issue.

1.5.2 Holography

Holography was first proposed as a data storage technology in the early 1960's, by a Polaroid scientist named Pieter J. van Heerden [39]. Ten years later, RCA Laboratory's scientists recorded 500 holograms in a specially made crystal, and 550 holograms in a light sensitive polymer material [1]. Even with those great results, due to the absence of cheap and appropriate components and the crescent successful researches on magnetic and semiconductor memories, the development of holographic data storage artifacts was placed on hold.

Nevertheless, the recent advent of appropriate components (such as more powerful laser diodes) and the low prices achieved by the needed material, finally yielded the retake of holographic data storage researches [16].

The latest published materials state that, with the current technologies, it is possible to develop an holographic storage system with a bandwidth of tens of GB/s, which is the optical limit. The problem with this system is its recording rate (only a few tens of KB/s). Nevertheless, the scientists say that future technologies might improve this value up to one GB/s [5].

Another interesting characteristic of holographic data storage systems is its density. Theoretically it could reach a density of tens of Tb/cm³, but, there is no such precise technologies yet to allow this density. Recent experiments can only record some tens of Mb/cm³, but some researchers say they would reach a density of one or two Gb/cm³ in a short time, which can lower its cost [5].

An example of a Holographic System is showed in Figure 1.6. The SLM (Spatial Light Modulator) contains the data that will be recorded. This processes is done by the interference generated by the laser beam that illuminates de SLM and the called reference laser beam, which comes from the Reference Array. To read data, only the Laser beam illuminates the fully opened SLM, and the image of the wanted page is projected on the Detector.

There are many multinational companies that have research teams working on Holography, some of those are IBM [11], Lucent [20], Kodak [4], Imation [12], Panasonic [22], Sony [36] and NEC [27]. There are also other independent companies working in this area. Although the holography can reach a bandwidth of 53 GB/s, due to the physical limits of the materials currently used in the computer industry (up to 10 GB/s) this value can not be effectively reached, once there is the need for an silicon made interface between the optical device and the commonly used computers [5]. The unique known solution to this problem is the development of a totally optical computer, which may take some years to take place.

1.5.3 Magneto-resistive Random Access Memory (MRAM)

Magnetism has being used as data storage technology since the end of the XIX century, when the Danish scientist Valdemar Poulsen invented the Telegraphone [34]. After that, many other scientists around the world proposed and implemented several improvements for this technology. Following those




Figure 1.6: An example of holographic system.

ideas, in 1956 the IBM Corporation [11] developed a magnetic system that stored 5 Mb of data in 50 double-sided disks, each one with a diameter of two feet. Nowadays, using nanotechnology and magneto-resistive inventions (e.g. magneto-resistive heads), researchers have achieved a density of a few tens of Gbits/in². But because of the nature of mechanical components, the seek time is still probably big.

The MRAM (Magneto-resistive Random Access Memory) idea was conceived through the use of high density of the ferromagnetic materials and its access time improvement. MRAM stores information in magnetic mode instead of electronic, as in DRAM (Dynamic Random Access Memory), and the orientation of magnetization is inside a thin ferromagnetic film. Its data is stored for long time periods without the need for external power supplies, like battery-backed CMOS. MRAM is considered a nonvolatile technology [40] because its physical state is modified when data is stored.

There are a lot of important companies working together to develop a commercial version of MRAM, but the actual prototypes are still small. The largest prototype presentation was made by Motorola Inc. in June, on the 2002 VLSI Symposium on Technology and Circuits. The device was a 1 Mbit MRAM, and it was also the first component to be integrated with the CMOS technology. In the same symposium, the Sony Corporation also presented a MRAM chip, smaller than the Motorola's chip [18].

As said above, this technology is still a laboratory product, and there is little information about it. Nowadays the most acceptable conclusion is that it is a very interesting technology, and theoretically, it is the closest technology to be the future "perfect memory semiconductor", because it is fast, non-volatile, and because of its low power consumption and unlimited read/write memory. Some announcements by Motorola, the first commercial versions of this memory will be on the market in 2004.

1.5.4 Nanotechnology

IBM developed a way to use their *Atomic Force Microscope* [AFM] in storage systems called *AFM Thermo-mechanical Data Storage*. The first prototype, *Millipede*, is already done and the results are very promising.

IBM began the research of this system in the early '90s. The initial model has a polymer disk and one cantilever whose density goes up to 30 Gb/s. Currently, it is used as an array of cantilevers that moves over a polymer medium. The write/over-write cycle reach 100.000 times.

The Millipede prototype have an array of 32x32 cantilevers acting over a polymer surface of 3x3 mm and the density reaching up to 400 Gb/in². A model with only one cantilever could have up to 1 Tb/in², but it loses read/write quality. Currently, data rates are low (few kilobytes per second per tip) but studies shows that this rate can be increased (some gigabits per second for each tip).

1.5.5 Chalcogenide Random Access Memory [C-RAM]

Initially developed by Energy Conversion Devices [ECD], exclusive licensed to Ovonyx, this memory is also called *Ovonic Unified Memory [OUM]*. It stores bits of information by changing the material structure state to crystalline or amorphous, in the same principle of CDR/CDRW/DVDRAM but using electrical charges instead of laser beam to write (this effect is called *phase-change*).

In amorphous state the material shows high-resistance and non-reflective characteristics, but in crystalline form the material shows low-resistance and reflectiveness. In the case of optical storage this behavior is used to control the laser beam reflection, in C-RAM it is used to control the resistance of the material.

Advantages over flash memories are noticeable in the actual stage: overwrite and bit-addressable capabilities and greater write-cycle (10^{13}) and write performance, projections demonstrate a potential speed up to DRAM performance. Intel and ST Semiconductors already signed development partnership with Ovonyx.

1.5.6 Other Technologies

The following technologies are those that are too experimental or still needing more technological advances to get competitive. There are also difficulties in getting information about them.

FRAM (Ferroelectric Random Access Memory): this technology is patented and developed by Ramtron International Corporation [31], and is being commercialized since 1992. Although it has been on the market for more than one decade, the biggest memory has a size of 256 Kb, which demonstrates a slow advance on this technology. Nevertheless it is a non-volatile, fast read/write and lower on power consumption, but has a destructive read, which imposes a limit to its read and write cycles. Its read/write time is much bigger than that of other memories, but its power waste to write and read is really low (estimated at 1nJ to write/read 32 bits, while other memories spend some microJ to make it) [35].

Figure 1.7 shows an example of what this technology looks like. As showed, the electric field moves the center atom (blue) to designate its state (0 or 1).




Figure 1.7: FeRAM example.

Polymer Memory: this technology is based on polymer materials with special formulation. This

memory is non-volatile, non-destructive read, high density, low on voltage and low on power consumption. Another interesting feature of this kind of memory is its low cost. There is no estimate to when this technology will be on the market, since it is still a laboratory product.

PFRAM (Polymeric Ferroelectric Random Access Memory): this memory store data by changing the polarization of polymer, which is between metal lines. This is a non-volatile memory and has fast read and write speed (initial read and write speed comparable to flash). The problem with this technology is its destructive read. Intel [15] is one example of companies that are investing on this technology [21].

Chapter 2

Operating Systems and Flash Memory

Flash memory has been accepted as the industry “de facto” standard solution for non-volatile storage in embedded systems. However, turning a “raw” flash memory chip into a usable disk-replacement for embedded applications is a rather complex task. This chapter discusses the use of flash memory in embedded systems from the point of view of the operating system, including device drivers, filesystems, and update support.

2.1 Device Drivers

Some flash memory gadgets include additional components that emulate an ordinary hard disk to the operating system. Such devices are usually interfaced to a filesystem through a traditional hard disk device driver that knows nothing about the peculiarities of flash memory operation. Nevertheless, ordinary flash memory units, as used in most systems, do not include such additional disk-emulation logic. They rely on specially designed device drivers to interface them to the file subsystem’s disk-like interface.

A *flash memory device driver* differs from a RAM-disk driver in that, besides mapping volume blocks into memory pages, it must also care for an efficient management of flash memory limitations like sector lifetime (for erasing), intra-sector rewriting of data, etc. Some device drivers will handle flash particularities internally, exporting a disk-like interface that supports the installation of an ordinary file system on a flash memory device. A second approach would be to export an interface with some flash-specific services. This second approach could probably lead to a more effective flash usage, but it would also require a flash-specific filesystem. Therefore, most systems opt for flash device drivers that emulate an ordinary disk.

2.1.1 Flash-specific Tasks

Independently of the strategy chosen to export services, a flash driver must handle several flash-specific tasks. The most important are summarized in this section.

Traditional filesystems are designed to update data in place. The same disk sectors are constantly rewritten with new data. This is a reasonable mode of operation for magnetic media, but at odds with flash memory, since there is a limitation in the number of times a flash can be rewritten. Updating data in place on a flash would cause some sectors (e.g. directories nodes) to expire their lifetimes far earlier than some other sectors. Furthermore, today’s typical flash sectors (~ 64 Kbytes) are too large to be directly mapped as disk blocks (~ 512 bytes).

There are two basic strategies to support the update of a small disk block stored in a larger flash

memory sector: erase-before-write and remapping (not-in-place-update). The first strategy consists in coping the whole sector to a temporary sector, erasing it, coping back unaltered blocks, and writing the new contents of the block being updated to the proper offset in the just-erased sector. The temporary sector must be subsequently erased. The second strategy consists in having a translation table that maps logical disk blocks in physical portions of flash sectors (*frames*). In this way, updating a block can be achieved by writing the data into a new frame and updating the translation table. Afterward, the old frame is marked “dead for posterior clean-up procedures.”

The remapping strategy has obvious performance advantages over immediate rewriting. Besides, it allows for the homogeneous usage of sectors, so the whole flash ages uniformly (*wear-leveling*). However, it has an important shortcoming: the need for *garbage collection*. Flagging frames as “dead” means leaving garbage behind in flash sectors that must be later collected. Ideally, this operation would be delayed until all frames in a sector are flagged “dead”, so reclaiming the sector (erasing it and putting it back at the free frame/sector list) could be done without a single copy. In practice, however, it is probable that many sectors will contain a mixture of dead and live frames. In this case, live frames of a sector must be moved to other sectors before its dead frames can be reclaimed.

Confronting remapping advantages (performance and wear-leveling) with its disadvantages (garbage collection) usually leaves a positive result, specially if the garbage collection procedure is implemented in such a way that reclaiming is preventively performed in advance on flash’s idle time. Consequently, remapping became the most common approach to support data update in flash memories.

Other flash-specific duties of device drivers are:

Data protection: flash memory units often support write protecting individual sectors. This can be used to protect critical data such as bootstrap and operating system kernel.

Fault recovery: since most flash memory units do not implicitly validate write operations, a flash device driver must periodically check for data integrity. If a failure is detected, the driver can mark the corresponding sector as “bad” and try to rewrite the data in other place. Write failures can be an indicative of life-time expiration.

2.1.2 Case Studies

At present time, there are many *flash driver layers* available both for commercial and open-source operating systems. The most significant example of driver that supports windows-like filesystems is the proprietary *Flash Translation Layer*, while the *Memory Technology Device* driver is mostly adopted by open-source systems.

Flash Translation Layer (FTL): FTL is a sector-based flash manager that provides logical to physical sector mapping, thus enabling a flash to look like an ordinary disk to the operating system [13]. As a rule, FTL supports any ordinary filesystem to be installed on a flash, but it is mostly deployed with windows-like systems such as VFAT. In order to achieve this, FLT defines small “virtual blocks” that are dynamically mapped into the flash memory unit (not-in-place-update), thus granting wear-leveling.

As depicted in figure 2.1, FTL divides the flash into one or more *Erase Units* (EU). The size of an EU depends on the flash sector size. Each EU is divided in one *Erase Unit Header* (EUH), one *Block Allocation Map* (BAM) and several *Read/Write Blocks* (RWB). The EUH contains information about the Erase Unit such as its size and the size of RWBs. The BAM keeps allocation information about every RWB in the EU and is usually stored after EUH. It is arranged in 4-byte entries, each describing the state of a RWB (deleted, bad, free, or allocated).

Read/Write Blocks can store three different types of data: *Virtual Block Data* (VBD), *Virtual Block Map Pages* (VBM) and *Replacement Pages* (RP). VBD is the user-data block exported to the host filesystem. VBM contains information to perform address translation. It is organized as a table of 4-byte entries, each one pointing to the logical address on the flash where the corresponding VBD resides. The

virtual block number supplied by the host filesystem is used as an index into this table. RPs are hold recent updates to an associated VBM, thus extending its live-time.

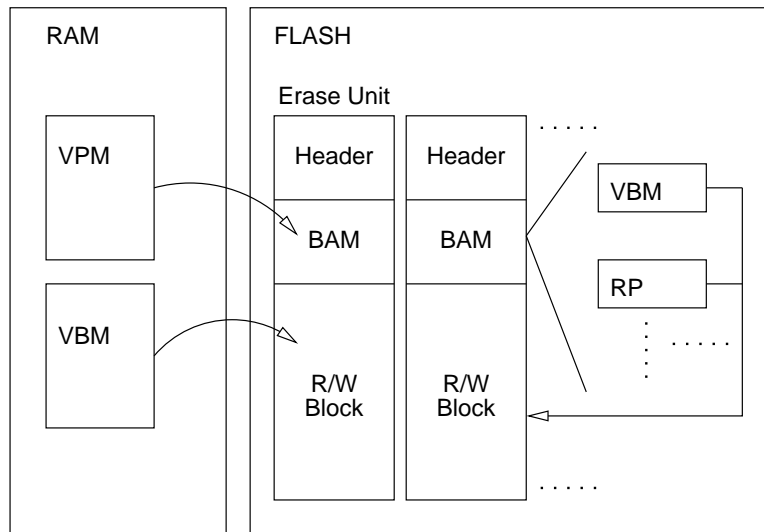


Figure 2.1: FTL architecture.

During initialization, FTL scans the Erase Unit Header and the Block Allocation Map of all Erase Units. Some flash sectors do not have an EUH and are taken as *spare blocks* or *transfer units* that are used on recovery/reclaim methods. Is not necessary to keep all the VBM stored on flash. Users can choose between keeping VBM in RAM, in flash or both of them. At initialization time, *Virtual Page Maps* (VPM) are created to hold the contents of those VMBs kept in flash (the same happening to Replacement Pages).

Memory Technology Device (MTD): MTD drivers [42] are a new class of drivers developed under Linux specifically for the embedded system area. The main advantage of MTD over conventional device drivers is that MTD drivers define a generic Linux subsystem for memory devices, in particular flash devices. Besides exporting the typical *raw block* interface for disk emulation, MTD drivers also export a *raw character* interface that allows filesystems to access the flash as if it were an ordinary linear memory device¹.

The main focus of the MTD project is to define a generic interface between hardware drivers and the upper layers of the operating system. In this way, hardware drivers need to know nothing about memory storage issues such as wear-leveling and garbage collection, simply providing routines for reading, writing and erasing. The MTD subsystem is divided in two types of modules: “users” and “drivers”. Drivers are the modules which provide raw read/write/erase access to physical memory devices. Users are the modules which use MTD drivers providing a higher-level interface to user-space. MTD subsystem architecture is depicted in figure 2.2.

2.2 Future Device Driver Technology:

According to the literature studied, the future of flash memory device drivers focus mainly on 2 ideas: (1) flexibility in the interaction between hardware and the system; and (2) system (re)adaption with existing and future management algorithms. Following this idea, HTD is the closest to that. Its capability of readapting to different algorithms and its flexibility in interacting with hardware and filesystem is the challenge to this project. This section discusses future drivers functionalities and its management.

¹The *raw character* interface is used by the Journaling Flash File System described later in section 2.3.

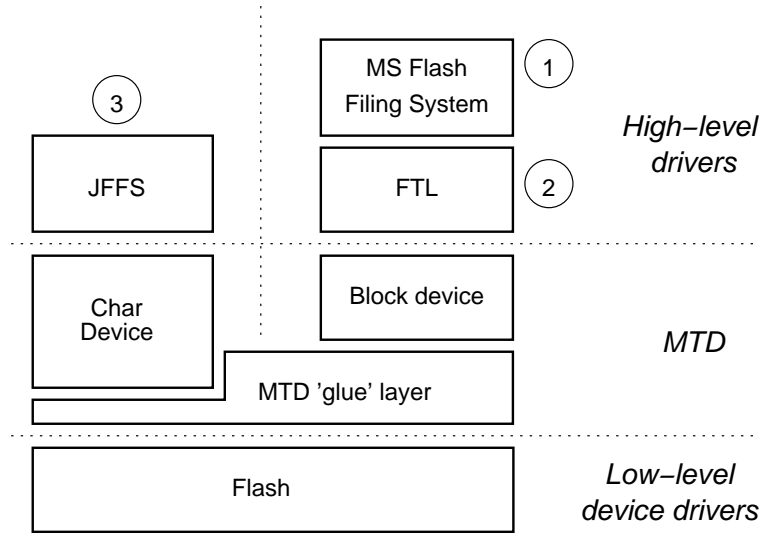


Figure 2.2: MTD Architecture.

Hardware Mediator Technology (HMT): HMT project consists of “low level software pieces” or “low level components” called pseudo-drivers, implemented for hardware devices specifically for the embedded system area. It has a “generic software philosophy” to interact with different kinds of hardwares like processors, buses, devices, controllers, etc. Hardware Mediators have been studied inside Application Oriented Object System (AOOS)[6] context by LISHA (Hardware and Software Integration Laboratory).

Initially, the HMT project focused on hardware mediator abstract elements for one hardware platform that would be used by EPOS[7] system abstractions and EPOS scenario aspects[8]. It was born principally to solve hardware specific architectural dependencies that implies a “layer overhead” and a “difficult software maintenance”, present in other “generic design technologies”.

However, these mediators are not intended to build a “universal virtual machine” for EPOS project, but hide the peculiarities of some hardware components. HMT used this idea to build a more efficient “pseudo-driver” layer specifically for RIFFS (showed later).

HMT, in its flash project context, differs from traditional drivers because they are compiled with an application (it means that HMT is just some code from an application point of view) but gives the idea of a driver layer for the software architecture designer.²

The aim of this “pseudo-driver” layer is similar to MTD (Memory Technology Device) where *hardware access layer* needs to know just basic operations like read, write and erase, ignoring the (complex) media management algorithms. The difference between these two technologies is the fact that HMT is not a “compiled-layer” to interact with hardware and filesystems.

As mentioned before, HMT is a group of components. Each HMT component implements one generic interface (read/write/erase, exported to file system) to interact with a specific piece of hardware, and its architecture is showed in figure 2.3.

When a filesystem needs to interact with the hardware, it does it via HMT, thus promoting the idea of transparent portability. Each hardware mediator itself is not portable, just its idea; they are specifically designed and implemented for each platform in an easy way (provided by its conception and its software engineer).

As mentioned before, mediators are themselves hardware dependent, being sometimes coded in assembly, using macros, or static metaprogramming techniques. Because HMT is compiled with its up-layers,

²HMT just exists in source code and it is never found “compiled alone”.

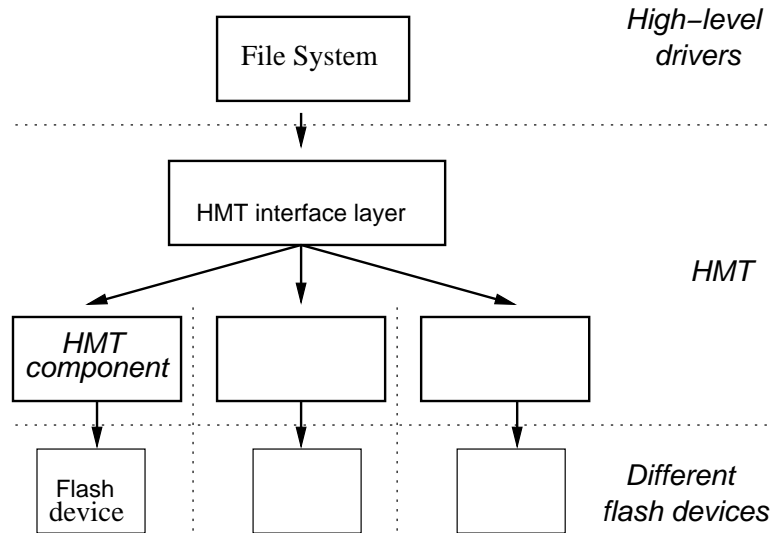


Figure 2.3: HMT Architecture.

different file systems will exist for each kind of flash memory, increasing performance compared to “generic software implementations”, because it doesn’t need to execute unnecessary code.

2.3 File Systems

Application programs are able to store and retrieve data from flash memory units through the services of a device driver. Very often, however, the approach of directly controlling the storage of data becomes inadequate, even in embedded systems. If different applications are to autonomously store data on the same flash memory, or if stored data is intensively manipulated, then installing a *filesystem* will usually be a more effective approach.

Although the filesystem scene is a very rich one, few filesystems have been proposed for flash memory. This is mainly due to the fact that most flash device drivers realize a disk-like interface that enables the installation of non-flash-specific filesystems on top of flash memory units. Nevertheless, a flash-specific filesystem will probably be advantageous in regard to a disk filesystem, since it has the opportunity to directly handle the limitations imposed by the technology.

2.3.1 Case Studies

True Flash File System (TrueFFS): M-Systems’ TrueFFS [23] implements its own patented standard called FTL, and exports the flash memory as a hard disk drive to the operating system. The FTL itself (implemented as a TrueFFS’s module) automatically handles wear-leveling, data protection, sector mapping, garbage collection, and fault recovery. Though called a file system, from a more strict point of view, TrueFFS would be considered a flash device driver. Indeed, TrueFFS requires the DOS FAT³ filesystem on top of it, and Figure 2.4 shows an example of how TrueFFS interacts with filesystem, operating system and FTL module. TrueFFS encapsulates the FTL module and then, exports block device services to filesystem and implements specific operating system integration.

Journalling Flash File System (JFFS): Axis Communications’ JFFS [2] differs from the traditional “virtual disk” approach (used by TrueFFS for example) while implementing a flash-specific filesystem

³DOS FAT is supported for all block devices




Figure 2.4: TrueFFS Layer inside Operate System context.

for embedded systems. JFFS version one (JFFS1) is a purely log-structured file system [33] and its improvement in version two [41] (JFFS2) includes compression and more efficient garbage collection.

The version one of JFFS (JFFS1), has two kind of structures: the (1) *node* structure (kept in flash device) called struct *jffs_raw_inode*; and the (2) *inode* struct (kept in RAM) is the struct as each node is associated. The inode header contains “the inode number”, “the inode data” and all “metadata” related with it (like the name of the inode and the parent’s inode number). Each node is associated with a version (hold in header) higher than all previous nodes belonging to the same inode number. This schematic is used by the host filesystem to know if the node is actually deleted. Also, there is a restriction on the maximum size of physical nodes, so large files will have many nodes⁴ associated with them.

JFFS1 has 3 main disadvantages: (1) compression is not supported; (2) hard links are also not supported; (3) inefficient garbage collection (because erase method get the first flash block in the log, without any criteria). An garbage collection inefficiency example is explained when the system contain a considerable amount of static data⁵: Suppose a system with 12MB of static data. In this system the garbage collection could move all of this static data.

The second version of JFFS (JFFS2) improves the disadvantages caused by JFFS1 and its code was intended to be portable, in particular to eCos, Red Hat’s embedded operation system [32]. While the original JFFS had only one kind of node, JFFS2 is more flexible⁶, allowing 3 types of nodes: (1) the *jffs2_nodetype_inode* is most similar to the struct *jffs_raw_inode* from version 1. It contains all the inode metadata as well as the range of data belonging to the inode. However, it no contains a file name nor the number of the parent inode. One important observation in this version 2 (as traditional UNIX-like filesystem), an inode is removed when the last directory entry referring to it has been unlinked, and there is no open file descriptors; (2) the *jffs2_nodetype_dirent* represents a directory entry, or a link to an inode. A link is removed by writing a dirent node with the same name but with target inode number zero - and obviously a higher version. With this difference between inodes and directory entries, the JFFS2 could support hard links; (3) the *jffs2_nodetype_cleanmarker* is written to a newly erased block to show that the erase operation has completed successfully and the block may safely be used for storage. This node type was introduced after JFFS2 had started to be used in real applications to avoid the problems of extensive power fail (because if power is lost during an erase operation, some bits in the medium may be left in an unstable state).

The second and third challenge in JFFS2 was its “data compression” and its “garbage collection improvement” respectively. Two compression algorithms were developed for use in JFFS2, and also the JFFS2 code can contain yet another copy of the *zlib compression library* which is present in the Linux

⁴Because symbolic links and especially device nodes, have small amounts of data, the data are not fragmented into many different nodes on the flash

⁵Static Data could be interpreted as libraries and executable files

⁶This flexibility is compared as traditional UNIX-like filesystems on each inode is a distinct entity from directory entries

Kernel source. Because the new garbage collection algorithm, the high-level layout of JFFS2 also changed from a single circular log format to an better performance decision collector format. This new format work with increased efficiency by collecting from one block at a time and making decisions about which block to garbage collect from next. To help garbage collection JFFS2 code also keeps a number of linked lists of structures representing individual erase blocks.

The majority of erase blocks will be on the *clean_list* or the *dirty_list*, which represent blocks full of valid nodes and blocks which contain at least one obsoleted node, respectively. In a new filesystem, many erase blocks may be on the *free_list*, and will contain only one valid node (a *jffs2_nodetype_cleanmarker*).

The operation of JFFS2 is very similar to that of the original JFFS where nodes, albeit now of various types, are written out sequentially until a block is filled, at which point a new block is taken from the *free_list* and writing continues from the beginning of the new block. Different of that three lists (*clean_list*, *dirty_list* and *free_list*) JFFS2 also keep in RAM a full map of filesystem. This map is constructed in mounting (by scanning the entire flash) and is supported by 2 structures: (1) *jffs2_inode_cache* and (2) *jffs2_raw_node_ref*. For each inode on the flash, there is a struct *jffs2_inode_cache* showed in figure 2.5 which stores its inode number, the number of current links to the inode, and a pointer to the start of a linked list of the physical nodes which belong to that inode⁷. Each physical node on the medium is represented by a smaller struct *jffs2_raw_node_ref*, also show in figure 2.5, which contains two pointers to other raw node references, the physical offset and total length of the node.

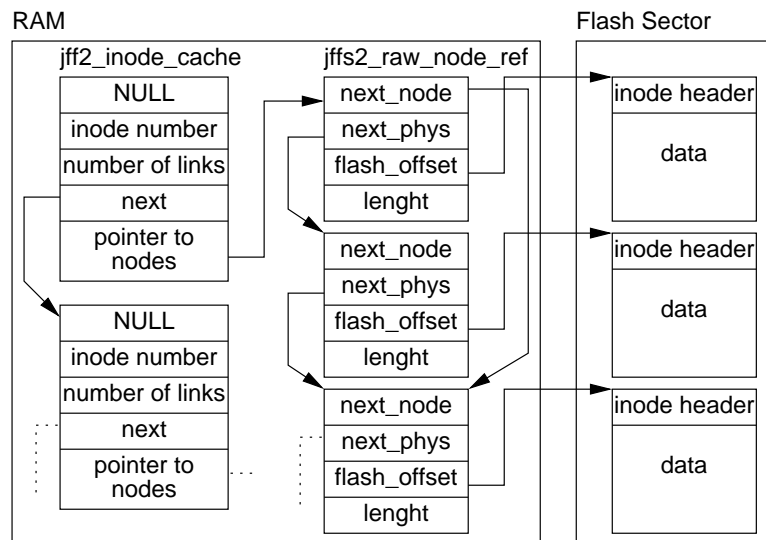


Figure 2.5: JFFS2 Architecture.

Mounting a JFFS2 filesystem involves a four-stage operation. (1) the physical medium is scanned and the *jffs2_inode_cache* are allocated and inserted into the hash table for each inode for which valid inodes are found and the *jffs2_raw_node_ref* either; (2) a first pass through all the physical nodes is made, obsoleted nodes can be detected and the “number of links field” is increased for each valid directory entry node; (3) a second pass is then made to find inodes which have no remaining links on the filesystem and delete them. (4) a third pass is made to free the temporary information.

Embedded File System (Efsys): *QNX Software Systems* Efsys [29] combines the functions of a filesystem manager and a device driver to manage a Flash filesystem on memory-based media. Because Efsys includes the device driver for controlling the hardware, there are separate versions of Efsys for different flash manufacturer. Each logical driver is called a *socket*, which consists of a contiguous and homogeneous region of flash memory. Each socket may be divided into one or more *partitions*. Actually there are two types of partitions supported: *raw partitions* and *flash filesystem partitions*. A raw partition is any partition that doesn’t contain a flash filesystem and may contain an image filesystem or some

⁷It is important to note that an inode is build by a linked list of nodes.

application-specific data. A flash filesystem partition exports the *POSIX-like* filesystem, which uses a QNX proprietary format to store data on the flash devices. This format stores files and directories as a direct linked list of *extents*. QNX defines *extent* like one consecutive region on one media (flash, disk, etc). Normally one file have multiples *extents*. Through literature studied, QNX flash filesystem supports the standard *POSIX* functionality like symbolic links, long filenames, etc; but there are one exception: it doesn't support "hard links". QNX flash filesystem supports transparently data decompression from flash. It is not true by compression ability, where the application need explicitly call a function to do that. When the flash is formatted by Efsys, the filesystem manager lock one (or more) untouchable block called *spare block*. This *spare block* will be used by garbage collection when system reclaim space. The QNX flash filesystem also has the expected features for flash filesystems like wear-leveling and fault recovery.

Other technologies:

- Virtual File System (VFS): PalmOS' VFS [28] requires the flash memory unit to be encapsulated in a device that emulates an ordinary storage device and therefore cannot be considered a flash filesystem. Nevertheless, it implements a sort of lightweight database system that may be of interest to some embedded applications. VFS stores both programs and data as collections of records that can be accessed via a database interface: a rather convenient approach for *Personal Information Manager* [PIM] applications.
- Microsoft Flash File System (MFFS): Microsoft's MFFS is a substratum for DOS FAT filesystems. Similarly to the TrueFFS, it handles all traditional flash tasks. One of its peculiarities is the adoption of variable size block, which can optimizes garbage collection: a large segment of garbage in a block can be converted into a new garbage-only block for immediate reuse.

2.4 Future File System Technology:

Since the beginning of flash filesystems, many algorithms have been suggested to overcome flash deficiencies. This section emphasizes on the last studies in flash management, that will become the next generation of this technology in a near future. Inside this new generation there is the management of index structures as showed in RIFFS and B-Tree Layer for Flash.

Reverse-Indirect Flash File System (RIFFS): *Research Group of Hardware and Software Integration Laboratory* (LISHA) has been studying a different manner to store data in (specifically) flash memories. The aim of the flash filesystem research project is to develop a complete flash filesystem totally flexible and customizable (generic in other words), but without compromising performance. To meet this expectation the group uses two non-usual ideas: (1) the rescue of the old paradigm *variable data-block length* and, (2) *reverse linked-list and indirect-tree management*, to compensate for flash memory disadvantages.

The purpose of RIFFS is not to develop a new garbage collection, nor a new wear leveling algorithm but implement a new way to store data in flash memories and try to group all the efficient flash management algorithms (or at least most of them) in an easy way and without overhead. Its structure is like a linked-list but indirect. In this manner, the filesystem tree appears in a reverse form, and this is the reason for its name. The RIFFS architecture is showed in figure 2.6.

With the reverse tree, RIFFS eliminates some aspects in garbage collection (like inside-block dirty pointers or logical to physical map), and *variable data-block length* improves wear leveling. To do its "job", there are two main data-structures: (1) *reverse_node*, present in media, and (2) *mounted_map_node*, present in RAM memory, as shown in figure 2.6.

A *reverse_node* just carries a logic number, its data and a pointer to a parent-node. On mounting, the entire flash is scanned and a direct-tree is built in RAM memory.

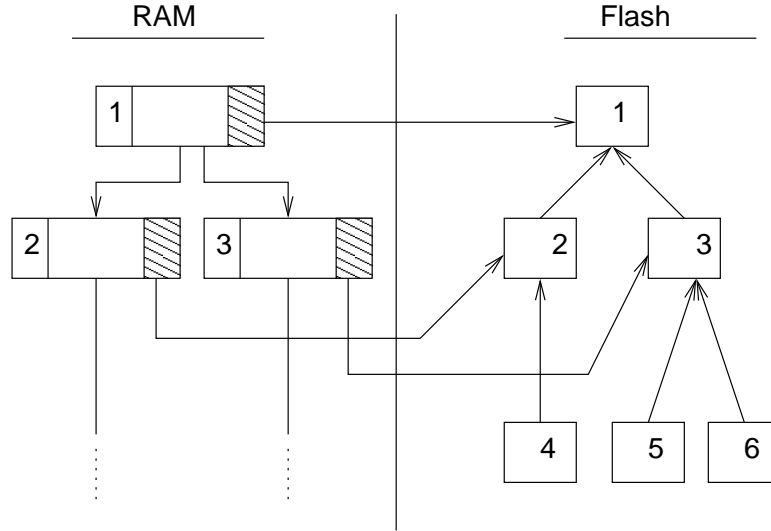


Figure 2.6: RIFFS Architecture.

B-Tree Flash Translation Layer (BFTL): Research Group of Computer Science Department of National Taiwan University proposed “An Efficient B-Tree Layer for Flash-Memory Storage Systems” [43].

For index structures (as B-Tree structures) which require major modifications, block-oriented access over flash memory could introduce a significant number of redundant writes (with data and pointers management, required by B-tree). This project proposes the efficient management of B-Tree index access over flash memory.

The focus of this project is on an integration of B-Tree index structures and the block-device emulation mechanism provided by FTL (Flash Translation Layer), resulting on an insertable module called BFTL. BFTL sits between the application layer and the block-device emulated by FTL. The BFTL module is dedicated to those applications which use services provided by B-Tree indexes. B-Tree index services requested by the upper-level applications are handled and translated by BFTL, and then block-device requests are sent from BFTL to FTL.

B-Tree node is physically represented by a collection of index units, and to help BFTL to identify the index units of the same B-Tree node, a *node translation table*(NTTab) is adopted. NTTab is introduced as an auxiliary data structure to collect index units efficiently. Basically, the *node translation table* maps a B-Tree node to a collection of index units supplied by FTL, and this index units could come from different sectors on flash. The *node translation table* is shown in Figure 2.7.

BFTL is constantly built in RAM memory, according to its node access. When a B-Tree node is visited, all the index units belonging to the visited node is collected and the NTTab could grow unexpectedly. To solve this problem researches propose to compact the *node translation layer*, and a variable (defined by user) is used to control the maximum length of the lists. After that, the compacted NTTab is written back to flash memory. This operation introduces an overhead in system, and obviously, the “system engineer” must choose between the compaction (in other words: the length of list) and the performance.

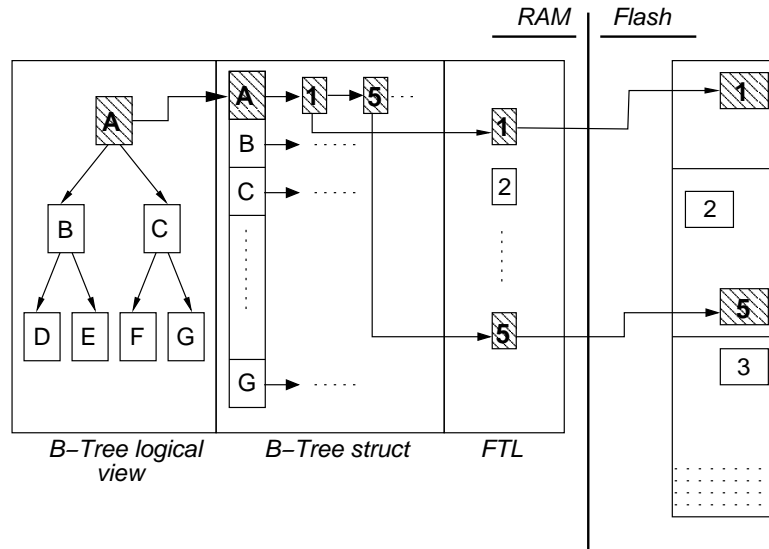


Figure 2.7: B-Tree node translation table.

2.5 Update Support

An apparently unimportant aspect of operating systems designed to operate from/on flash memories is the way they can be updated, for the time in which embedded systems were delivered with absolute, immutable firmware is long gone. Today’s embedded systems must consider regular upgrade operations for firmware, operating system, and applications — the popular “system re-flashing”.

One aspect in which the operating system can easy or difficult updating concerns the way it is stored in flash. If the system is a monolithic piece of software installed on a predefined location in flash, it is very likely the upgrading it will require the whole flash to be erased, since the installation of a larger system image could corrupt preexisting filesystem structures. Moreover, some systems are directly executed from flash and allow applications to make absolute references to internal functions. In this case, updating the system implies in updating all applications. Such a condition could have serious consequences for perpetual applications that store context information in the flash. For an example consider an *odometer* application on a car: re-flashing the system without preserving the current mileage count would transform an old chariot in a “0 Km”. Typical Windows CE installations present such shortcomings.

Re-flashing side-effects can be avoided, or at least minimized, if the operating system is itself installed in the filesystem. In this case, upgrading the operating system solely requires some files to be overwritten/added and has no direct effect on applications. Typical Linux distributions for embedded systems function in this way. In particular, the different Linux distribution provides a tool called “apt-get” that is able to download packages from the network (via “wget”) and install/upgrade them in a consistent way. Indeed, embedded Linux distributions usually inherit on-the-fly kernel upgrade mechanism from desktop distributions called *loadable kernel* modules.

2.5.1 Case Studies

WindowsCE uses the first approach so it have to rewrite the entire data in the memory. Pocket Linux otherwise uses the second approach and can even change kernel on-fly, updating just the selected modules (files).

WindowsCE for Automotive Update The current version of Windows CE does not include the necessary API to update all or part of the Flash. Fortunately, it was done by a small Operating System called *burnos* [24] (just in Automotive WindowsCE version [25]).

Burnos is a temporary second operating system that executes while the current OS is being reprogrammed. Because burnos reside in the same block as the boot code, if the update process is interrupted, the utility must be available to resume or redo the update. On power-up or reset, if the boot loader does not find a valid image in the rest of the Flash, it will invoke the update OS to install the program.

The WindowsCE operating system and applications are a single binary image, so they are written sequentially into the rest of the Flash without regard for block boundaries. To update the OS and the applications, The burnos performs some tasks showed on figure 2.8: (1) Take control of the system. Since Windows CE does not implement this capability or support a partial update, the utility must preempt the operating system; (2) Erase all blocks that do not contain boot code. The entire binary image must be erased and reprogrammed; (3) Input the new binary image. All code necessary to receive the new program must be included in the utility, because the rest of the code will have been erased; and (4) Program the new image into the Flash.

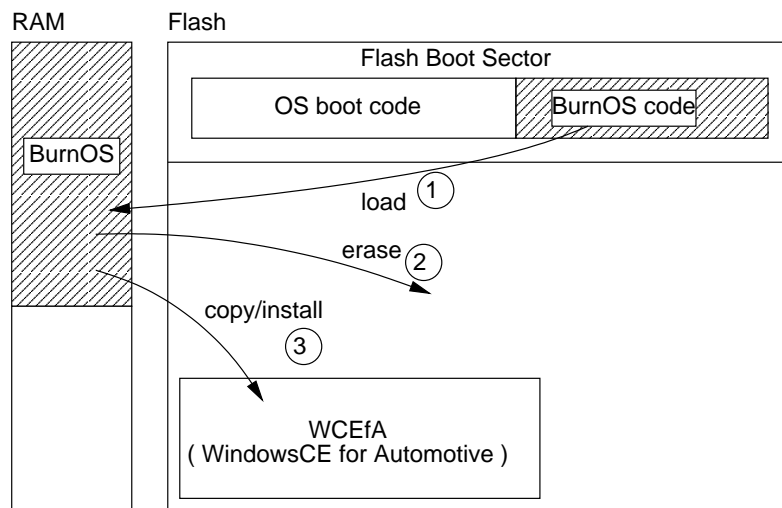


Figure 2.8: WCEfA update schematic.

During the 1st. phase, the system reboots 2 times. The initial boot happen just to clear RAM. Next, the system loads the burnos, and modify the system loader to boot burnos. The last 3 phases, there are three reboots: the first reboot happens to copy Burnos to RAM and then, it sets a special boot-mode and boots again. Next, the temporary OS started a configuration routine. Last, the update process loads the new image into RAM. Once loaded and verified, the flashing begins. After the flashing has completed, the boot-mode returns to original configuration and upon final reboot, the system comes up under the new version.

Linux Update support The most basic way to add functions (routines) in Linux Kernel is to put code inside it. But if there is necessity to divide it in distinct peaces (called modules), you can also add code by Loadable Kernel Module (usually called LKM) [19].

LKMs are attached to Linux Kernel dynamically by loading modules that implement protocol handlers, where these modules install function pointers into a protocol handler list. The LKM is divided in four kind of modules: (1) device drivers (used to communicate with a piece of hardware); (2) filesystem drivers (used to interprets the contents of a filesystem as files and directories); (3) system calls (used to get services from kernel) and (4) executable interpreters (used when an OS has more than one kind of executable files). The kernel isolates certain functions, including these, especially well so they don't have to be intricately wired into the rest of the kernel.

The most advantage, in updating context, is that LKM helps the system to update just modules, not the entire Kernel code, and this saves time to rebuild (recompiling) and installing it. The "module's load" can be done in three different ways: (1) automatically, using scripts; (2) manually from the command line (not usual in embedded systems); or (3) loaded directly from kernel as need (like devices drivers

for example), normally done in kernel boot. Linux uses itself tools to do this “job”, called modutils. Modutils is a package that has some commands to handle the modules (like load_module, unload_module, verify_module, etc).

An LKM lives in a single ELF object file, normally named like “file.o”. The “file.o” dependencies are resolved in load-time (dynamically). Although advantages of LKM, is common to find in embedded systems all the module’s code inside the kernel because its better performance ⁸. The figure 2.9 shows an example of how some library could be updated via LKM without recompiling/installing the entire system.

In phase one the (1)“lib.o” and (2)“startup.configuration” are erased; (3)“loaded.lib.o” is unloaded and (4)“kernel reference” is destructed. The second phase the (5)“new_lib.o” and (6)“new_startup.configuration” are record (7)“new_loaded.lib.o” is loaded and (8)“new kernel reference” are build. This update method is done without reboot the system. This method could vary for each Linux distribution.

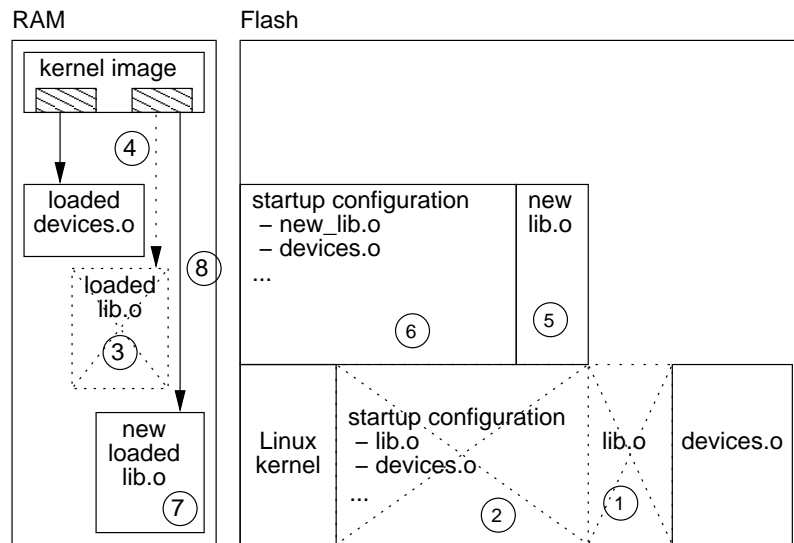


Figure 2.9: WCEfA update schematic.

In update context, there is also one important tool (or package) to do the system update and its name is “apt-get”. Apt-Get is a package management system that is designed on debian linux systems and ported to another distributions. It has the ability to connect to a server “repository” (on internet and/or locally) and checks (automatically) all “package dependencies” required. Because apt-get is implemented in GPL ⁹ [9] designers/programmers has the advantage to (re)adapt this software to its necessities. Apparently all the linux distributions has this capability to update dynamically modules and packages either in embedded systems.

⁸The balance between performance and flexibility can be defined by Software Architecture Engineer based in its knowledge.

⁹General Public License (GPL) is intended to guarantee your freedom to share and change free software—to make sure the software is free for all its users.

Bibliography

- [1] W. C. Stewart ET AL. Experimental read-write holographic memory. Technical Report RCA Review 34(1) 3-44, R.C.A., New York, N.Y., U.S.A., 1973.
- [2] Axis Communications. *The Journalling Flash File System (JFFS)*, online edition, December 2002. [<http://developer.axis.com/software/jffs/>].
- [3] Geoffrey W. Burr and Inmaculada Leyva. Multiplexed phase-conjugate holographic data storage with a buffer hologram. *OPTICS LETTERS*, 25(7), 2000.
- [4] Eastman Kodak Company. *Kodak*, online edition, March 2003. [<http://www.kodak.com>].
- [5] Jean-Jacques P. Drolet Ernest Chuang, Wenhai Liu and Demetri Psaltis. Holographic random access memory (hram). In *Proceedings of the IEEE, Vol. 87, No. 11*. IEEE, November 1999.
- [6] Antonio Augusto Frohlich. *Application-Oriented Operating Systems*. Number 17 in GMD Research Series. GMD - Forschungszentrum Informationstechnik, Sankt Augustin, August 2001.
- [7] Antonio Augusto Frohlich and Wolfgang Schröder-Preikschat. EPOS: an Object-Oriented Operating System. In *Proceedings of the 2nd ECOOP Workshop on Object-Orientation and Operating Systems*, pages 38-43, Lisbon, Portugal, June 1999.
- [8] Antonio Augusto Frohlich and Wolfgang Schröder-Preikschat. Scenario Adapters: Efficiently Adapting Components. In *Proceedings of the 4th World Multiconference on Systemics, Cybernetics and Informatics*, Orlando, U.S.A., July 2000.
- [9] GNU's Not Unix. *GNU General Public License*, online edition, December 2002. [<http://www.gnu.org/copyleft/gpl.html>].
- [10] Steve Grossman. Future trends in flash memories. In *Proceedings of the 1996 IEEE International Workshop on Memory Technology, Design and Testing (MTDT'96)*, Singapore, August 1996. IEEE.
- [11] IBM Corporation. *IBM Corporation*, online edition, March 2003. [<http://www.ibm.com>].
- [12] Imation Corporation. *Imation*, online edition, March 2003. [<http://www.imation.com>].
- [13] Intel Co. *Understanding the Flash Translation Layer (FTL) Specification*, online edition, December 1998. [<ftp://download.intel.com/design/flcomp/applnots/29781602.pdf>].
- [14] Intel Co. *What is Flash Memory*, online edition, October 2002. [<http://www.intel.com/design/flcomp/articles/298312.pdf>].
- [15] Intel Corporation. *Intel Corporation*, online edition, March 2003. [<http://www.intel.com>].
- [16] J. Ashley, M. P. Bernal, G. W. Burr, H. Coufal, H. Guenther, J. A. Hoffnagle, C. M. Jefferson, B. Marcus, R. M. Macfarlane, R. M. Shelby, and G. T. Sincerbox. Holographic data storage. *IBM Journal Of Research And Development*, 44(3), 2000.
- [17] R. Katti. Solid State Memory Study Final Report. Technical Report JPL-1994-539, Jet Propulsion Laboratory, St. Paul, U.S.A., February 1994.

- [18] David Lammers. The promise of mram. *EE Times*, July 2002. This article relates the presentations of 2 MRAM chips at 2002 Symposium on VLSI Circuits, one made by Motorola, this chip is a 1 Mbit MRAM chip, integrated with the CMOS tech, and other 256 Kbit chip made by Sony. It also verses about the IBM's opinion on this technology.
- [19] The Linux Documentation Project. *Loadable Kernel Modules*, online edition, December 2002. [<http://www.tldp.org/HOWTO/Module-HOWTO/>].
- [20] Lucent Technologies. *Lucent Technologies - Bell Labs Innovations*, online edition, March 2003. [<http://www.lucent.com>].
- [21] Mike Magee. Intel to intro 'internet on chip' this year. *The Inquirer*, February 2002. This article relates the Intel's announcements about its PFRAM chips, and its intentions to produce this kind of memory chips in large scale. It also relates that Intel is evaluating other non volatile technologies, such as Ovonic Unified Memory (OUM), MRAM, FeRAM.
- [22] Matsushita Electric Industrial Co., Ltd. *Panasonic*, online edition, March 2003. [<http://www.panasonic.com>].
- [23] Microsoft Co. *Linear Flash Memory Devices on Microsoft Windows CE 2.1*, online edition, December 2002. [<http://msdn.microsoft.com/library/default.asp?url=/library/en-us/dnce21/html/flash21.asp>].
- [24] Microsoft Windows. *OS Update Support*, online edition, December 2002. [<http://msdn.microsoft.com/library/>].
- [25] Microsoft Windows. *Windows CE for Automotive*, online edition, December 2002. [<http://www.microsoft.com/automotive/winceauto/>].
- [26] Inc. Motorola. Motorola sets major milestone with 1 mbit mram universal memory chip with copper interconnects, July 2002. This is an Motorola's press release about the presentation of a 1 Mbit memory chip at the 2002 VLSI Symposia on Technology and Circuits.
- [27] NEC Corporation. *NEC*, online edition, March 2003. [<http://www.nec.com>].
- [28] PalmSource, Inc. *Virtual File System*, online edition, December 2002. [<http://www.palmos.com/dev/support/docs/>].
- [29] QNX Neutrino OS. *QNX Momentics Non-Commercial Documentation Roadmap*, online edition, December 2002. [http://www.qnx.com/developer/docs/momentics_nc_docs/roadmap/].
- [30] Kamal Rajkanan. Flash memory technologies — a review. In *Proceedings of the 1996 IEEE - International Workshop on Memory Technology, Design and Testing (MTDT'96)*, Singapore, August 1996. IEEE.
- [31] Ramtron International Corporation. *Ramton*, online edition, March 2003. [<http://www.ramton.com>].
- [32] Red Hat, Inc. *ECOS - Embedded Configurable Operating System*, online edition, December 2002. [<http://sources.redhat.com/ecos/>].
- [33] Mendel Rosenblum and John K. Ousterhout. The Design and Implementation of a Log-structured File System. *ACM Transactions on Computer Systems*, 10(1):26–52, 1992.
- [34] George Saliba. Laser guided magnetic recording. *Computer Technology Review*, XX(7), 2000.
- [35] Ali Sheikholeslami and P. Glenn Gulak. A survey of circuit innovations in ferroelectric random-access memories. *Proceedings Of The IEEE*, 88(5):667–689, jul 2000.
- [36] Sony Corporation of America. *Sony*, online edition, March 2003. [<http://www.sony.com>].
- [37] ST Microelectronics. *Background Information: Flash Memories*, online edition, February 2001. [<http://us.st.com/stonline/press/news/back2002/b979m.htm>].

- [38] D. A. Thompson and J. S. Best. The future of magnetic data storage technology. *IBM Journal Of Research And Development*, 44(3), 2000.
- [39] P. J. van Heerden. A New Optical Method Of Storing And Retrieving Infrmation. *Appl. Opt.*, 2(4):387–392, 1963.
- [40] Frank Wang. A modified architecture for high-density mram. *ACM SIGARCH Computer Architecture News*, 29(1):16–22, 2001.
- [41] David Woodhouse. JFFS: The Journalling Flash File System. In *Proceedings of the Ottawa Linux Symposium 2001*, Ottawa, Canada, July 2001.
- [42] David Woodhouse. *Memory Technology Device*, online edition, December 2002. [<http://www.linux-mtd.infradead.org/tech/>].
- [43] Chin-Hsien Wu, Li-Pin Chang, and Tei-Wei Kuo. An efficient b-tree layer for flash-memory storage systems. *The 9th International Conference on Real-Time Computing Systems and Applications (RTCSA)*, 2003.