AN APPLICATION-ORIENTED APPROACH FOR THE GENERATION OF SOC-BASED EMBEDDED SYSTEMS

Fauze V. Polpeta and Antônio A. Fröhlich
Federal University of Santa Catarina, Brazil
{fauze,guto}@lisha.ufsc.br

Abstract

This paper outlines a strategy for automating the design of embedded systems including their hardware and software components. We focus in the *Hardware Mediator* construct, a portability artifact that was originally proposed to enable the port of component-based operating systems to very distinct architectures. Besides giving rise to a highly adaptable system-hardware interface, these mediators are approached as a new co-design artifact that can be used to enable the generation of customized *system-on-a-chip* instances and the associated run-time support systems considering the requirements of target applications.

Keywords: Application-Oriented System Design, Embedded Systems, Hardware Mediators, Operating Systems, System-on-a-Chip

1. INTRODUCTION

Embedded systems are becoming more and more complex, yet, there is no room for development strategies that incur in extended time-to-market in this extremely competitive sector. In this context, the *System-on-a-Chip* (SoC) define a compromise between system complexity and development costs [2]. Furthermore, the advances in programmable logic devices (PLD) are enabling developers to instantiate and to evaluate complex SoC designs in a short period of time. This can drastically decreases the time-to-market and turns PLDs an important technologic alternative in the development of embedded systems.

Indeed, some *embedded systems* can be completely implemented in hardware using the SoC approach. However, the more complex the application, the greater is the probability it will need some kind of *run-time support system* and an *application program*. This is, after all, the reason why so many groups are concentrating efforts to develop processor soft cores such as Leon2 and OpenRisc [11]. Nevertheless, run-time support systems are often neglected by

currently available SOC development methodologies and tools, being mostly restricted to simple processor scheduling routines and the definition of a hardware abstraction layer. The gap between software and hardware gets even bigger when we recall that one of the primary goals of an operating system is to grant the portability of applications, since ordinary operating systems cannot go with the dynamism of SOCs.

In this paper we discuss the use of *Hardware Mediators* [16] to enable the automatic generation of SoC-based embedded systems. The deployment of *Application-Oriented System Design* (AOSD) [4] on the context where hardware mediators were originally proposed—*software-hardware interfacing*—fosters this portability artifact to a new perspective on the design of embedded systems. Mediators are figured as pointers for generating a "machine description" that matches, in association with a run-time support system, the requirements of dedicated applications. The following sections describe the basics of the AOSD method, the concepts of hardware mediators and how these mediators can be deployed on the generation of SoC-based embedded systems. Subsequently, in a experimental case study, we consider the EPOS system [5], an application-oriented operating system that relies on hardware mediators to foster portability and also to enable automatic hardware generation. The paper is closed with a discussion of related works and the author's perspectives.

2.   APPLICATION-ORIENTED SYSTEM DESIGN

Application-Oriented System Design (AOSD) [4] proposes some alternatives to proceed the engineering of a domain towards software components. In principle, an application-oriented decomposition of the problem domain can be obtained following the guidelines of *Object-Oriented Decomposition* [3]. However, some subtle yet important differences must be considered. First, object-oriented decomposition gathers objects with similar behavior in class hierarchies by applying variability analysis to identify how one entity specializes the other. Besides leading to the famous "fragile base class" problem [12], this policy assumes that specializations of an abstraction (i.e. *subclasses*) are only deployed in presence of their more generic versions (i.e. *superclasses*).

Applying variability analysis in the sense of *Family-Based Design* [15] to produce independently deployable abstractions, modeled as members of a family, can avoid this restriction and improve on application-orientation. Certainly, some family members will still be modeled as specializations of others, as in *Incremental System Design* [7], but this is no longer an imperative rule. For example, instead of modeling connection-oriented as a specialization of connectionless communication (or vice-versa), what would misuse a network that natively operates in the opposite mode, one could model both as autonomous members of a family.

A second important difference between application-oriented and object- oriented decomposition concerns environmental dependencies. Variability analysis, as carried out in object-oriented decomposition, does not emphasizes the differentiation of variations that belong to the essence of an abstraction from those that emanate from the execution scenarios being considered for it. Abstractions that incorporate environmental dependencies have a smaller chance of being reused in new scenarios, and, given that an application-oriented operating system will be confronted with a new scenario virtually every time a new application is defined, allowing such dependencies could severely hamper the system.

Nevertheless, one can reduce such dependencies by applying the key concept of *Aspect-Oriented Programming* [8], i.e. aspect separation, to the decomposition process. By doing so, one can tell variations that will shape new family members from those that will yield scenario aspects. For example, instead of modeling a new member for a family of communication mechanisms that is able to operate in the presence of multiple threads, one could model multithreading as a scenario aspect that, when activated, would lock the communication mechanism (or some of its operations) in a critical section.

Based on these premises, Application-Oriented Systems Design guides a domain engineering procedure (see Figure 1) that models software components with the aid of three major constructs: families of scenario-independent abstractions, scenario adapters and inflated interfaces.

Families of scenario independent abstractions

During domain decomposition, abstractions are identified from domain entities and grouped in families according to their commonalities. Yet during this phase, aspect separation is used to shape scenario-independent abstractions, thus enabling them to be reused in a variety of scenarios. These abstractions are subsequently implemented to give rise to the actual software components.

Scenario adapters

As explained earlier, AOSD dictates that scenario dependencies must be factored out as *aspects*, thus keeping abstractions scenario-independent. However, for this strategy to work, means must be provided to apply factored aspects to abstractions in a transparent way. The traditional approach to do this would be deploying an *aspect weaver*, though the *scenario adapter* construct [6] has the same potentialities without requiring an external tool. A scenario adapter wraps an abstraction, intermediating its communication with scenario-dependent clients to perform the necessary scenario adaptations.
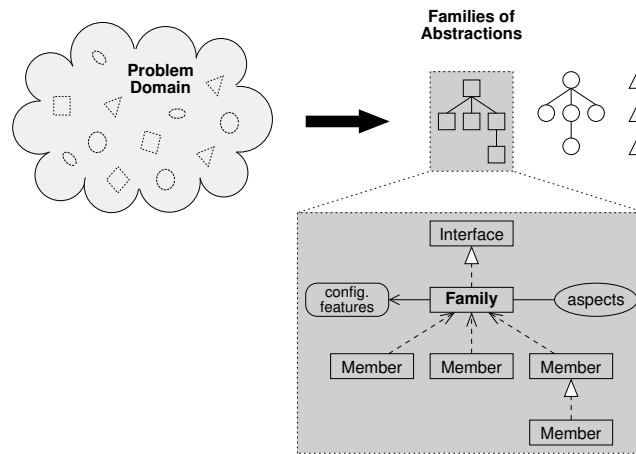
Figure 1.   Overview of application-oriented domain decomposition.

Inflated interfaces

Inflated interfaces summarize the features of all members of a family, creating a unique view of the family as a "super component". It allows application programmers to write their applications based on well-know, comprehensive interfaces, postponing the decision about which member of the family shall be used until enough configuration knowledge is acquired. The binding of an inflated interface to one of the members of a family can thus be made by automatic configuration tools that identify which features of the family were used in order to choose the simplest realization that implements the requested interface subset at compile-time.

3.   HARDWARE MEDIATORS

An operating system designed according to the premises of Application-Oriented System Design can be summarily viewed as sets of software components that can be configured, adapted and integrated in order to give rise to highly customized and scenario-specific instances of run-time support systems. However, besides all the benefits claimed by software component engineering, such a class of run-time support systems is prone to the same need for portability as their more conventional relatives.

Traditional portability strategies, mainly focused in hardware abstraction layers (HAL) and virtual machines (VM), are not concerned with the AOSD's purposes. Being a product of a system engineering process (instead of a domain engineering process), these strategies usually build a monolithic abstraction layer that encapsulates all the resources available in the hardware platform without properly regarding the application needs. Such modeling may

be a problem when the platforms to be interfaced are based on SoCs. The diversity of architectures and devices in these platforms lead us to diagnose that the traditional specification techniques for sw-hw interfacing are still far from the ideal "plug-and-play" [14]. In addition, whenever SoCs are built on *Programmable Logic Devices* such as FPGAs, the hardware specifications can be modified in a short period of time [17], and thus compromising much more the portability of the system.

In order to deal with this dynamism and to foster the portability of system abstractions to virtually any architecture, a system designed according to the concepts of AOSD relies on the hardware mediator construct. As discussed in *Hardware Mediators: a Portability Artifact for Component-Based Systems* [16], the main idea behind this portability artifact is not to build an universal hardware abstraction layer or virtual machine, but sustaining an *interface contract* between the operating system and the hardware. Each hardware component is mediated via its own mediator, thus granting the portability of abstractions that use it without creating unnecessary dependencies. Indeed, hardware mediators are intended to be mostly static-metaprograms and thus "dissolve" themselves in the system abstractions as soon as the interface contract is met. In other words, a hardware mediator delivers the functionality of the corresponding hardware component through a system-oriented interface.

An important element of hardware mediators are *configurable features*, which designate features of mediators that can be switched on and off according to the requirements dictated by abstractions. A configurable feature is not restricted to a flag indicating whether a preexisting hardware feature must be activated or not. Usually, it also incorporates a *Generic Programmed* [13] implementation of the algorithms and data structures that are necessary to implement that feature when the hardware itself does not provide it. An example of configurable feature is the generation of CRC codes in mediators that abstract communication devices.

Likewise abstractions in AOSD (Figure 1), hardware mediators are organized in families whose members represent significant entities in the hardware domain. Such modeling enables the generation of object-code only for those mediators that are necessary to support the application. Non-functional aspects and cross-cutting properties are factored out as *scenario aspects* that can be applied to family members as required. For instance, families like UART and NIC (*Network Interface Card*) must often operate in exclusive-access mode. This could be achieved by applying a share-control aspect to the families.

4.   CO-DESIGNING WITH HARDWARE MEDIATORS

Although originally devised to give rise to highly adaptable system- hardware interface, hardware mediators can be also used for generating application-

oriented hardware instances. More specifically, in the context of programmable logic, where hardware components, namely soft-IPs, are described using hardware description languages (e.g. VHDL, VERILOG) in order to implement the elements of the underling hardware technology. Hence, by associating hardware mediators with those descriptions one can infer which hardware components are necessary to support the application.

A component-based system can thus rely on hardware mediators not only to interface its system abstractions to the hardware, but also to dictate which components will build-up the target hardware. As soon as a hardware mediator is selected to interface a hardware component, the associated IP is selected from a repository in order to integrate a hardware description that will be synthesized in a PLD device. Such a description in the form of a SoC would embed only the hardware components necessary to support the run-time system and, in turn, the application.

For instance, consider a family of `NIC` mediators and a family of `NIC` hardware components whose members are associated to the members of the mediator's family. From a given application that uses a system abstraction to implement an Ethernet network one can infer that a member of the family `NIC` will be instantiated. However, the decision of which specific member will be instantiated, since all members are functionally equivalent, is up to the application's programmer. This situation characterizes what we named a *combined IP-selection*. The hardware devices are inferred considering an application requirement and a specific decision of the application's programmer.

Another scenario, named *discreet IP-selection*, is related to the selection of hardware components considering only the application's requirements — no explicit programmer decision must be taken. A good example for this scenario is related to the memory management scheme that will be implemented by the run-time support system. Once the application programmer uses system abstractions that rely on a *paging* scheme (e.g. *multitasking*), the `MMU` mediator is automatically inferred and, consequently, a memory management unit will be selected for synthesis. Conversely, when a *flat* memory scheme is adopted no memory management unit is synthesized.

A third scenario, named *explicit IP-selection*, represents the chance of the programmer to choose the hardware components that will be instantiated in the system. Even if the respective mediators are "hidden" by system abstractions, the programmer explicitly selects the hardware components that he is intending to embed in the SoC. Indeed, this selection strategy is always taken when the programmer initially specifies which architecture model (e.g. `SPARCV8`, `OR32`) the system will follow.

Furthermore, the approach is not restricted to the specification of which IPs will be instantiated in a PLD device. The element *configurable feature* explained earlier can be also deployed on hardware components. They can be

used to configure IPs in order to properly support the application and also to switch on and off some functionalities that can be implemented in hardware or in software. As exemplified earlier with CRC codes, a configurable feature can be used to enable the generation of these codes in a `UART` mediator for data transmitted over a serial communication line. Such codes, otherwise, could be generated by the hardware itself instead of the software. In this case, since the IP that implements the `UART` device supports generation of CRC codes, a hardware configurable feature would be activated in order to enable the synthesis of these functionalities in the SoC.

5.   CASE STUDY: SOCS IN EPOS

The Embedded Parallel Operating System (EPOS) aims at delivering adequate run-time support for dedicated computing applications. Relying on the *Application-Oriented System Design* method, EPOS consists of families of software components that can be adapted to fulfill the requirements of particular applications. In order to maintain the portability of its systems abstractions and to enable the generation of application-oriented SoCs, the EPOS system relies on the hardware mediator construct.

An application written on EPOS can be submitted to a tool that scans it searching for references to the *inflated interfaces*, thus rendering the features of each family that are necessary to support the application at run-time. This task is accomplished by a tool, the analyzer, that outputs an specification of requirements in the form of partial component interface declarations, including methods, types and constants that were used by the application.

The primary specification produced by the analyzer is subsequently fed into a second tool, the configurator, that consults a build-up database to create the description of the system's configuration. This database holds information about each component in the repository, as well as dependencies and composition rules that are used by the configurator to build a dependency three. The output of the configurator consists of a set of keys that define the binding of inflated interfaces to abstractions and activate the scenario aspects eventually identified as necessary to satisfy the constraints dictated by the target application or by the configured execution scenario. On the side of the hardware components, the configurator produces a list of instantiated mediators and specifies which of these mediators promote IP synthesis.

The last step in the generation process is accomplished by the generator. This tool translates the keys produced by the configurator into parameters for a statically metaprogramed component framework and causes the compilation of a tailored operating system instance. In addition, whenever a SoC needs to be tailored, the `generator`, based on the IP configurable features, produces a synthesis configuration file. This file, as well as the selected IPs, are passed to
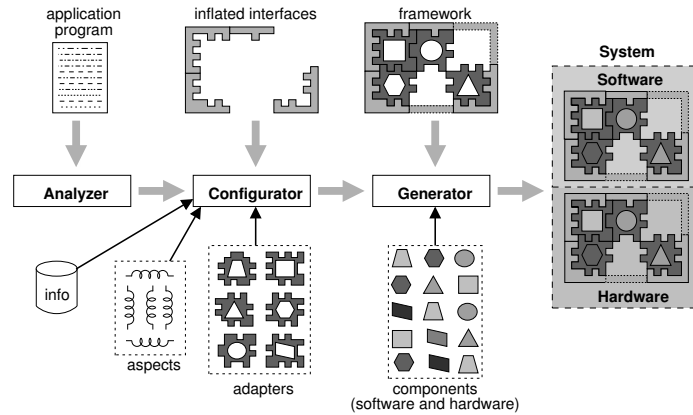
Figure 2.   An overview of the tools involved in the automatic generation of run-time support and hardware instances.

a third-party tool, which in turn performs the synthesis of a SOC. An overview of the whole procedure is depicted in Figure 2.

5.1   Sample Instance: Leon SoCs in EPOS

In order to evaluate our approach for automating the design of SoC-based embedded systems we used the *Leon2 Processor*. This "processor" was created in order to enable the development of customized SoCs based on the SPARCV8 CPU core. The "modular" design of LEON2 enable us to specify which of its IPs will be synthesized in the SOC. The logic necessary to glue the IPs is implicitly defined in the source-code through a set of programming asserts, which are, in turn, used to properly configure and plug-in the IPs in a AMBA bus framework [1]. The Figure 3(a) shows the block diagram of LEON2. Besides the CPU core, LEON2 includes a set of peripherals that can be plugged in as soon as the user selects them.

The experiments were performed on a Xilinx Virtex2 FPGA development kit and consisted of evaluating an application for which a run-time support system and a SOC should be generated. The application implemented two threads, *TX* and *RX*, which were executed in a cooperative environment in order to send and receive data through an UART. Aiming at signaling the *RX* thread to deal with new data in the UART buffer the mechanism of *interrupts* was used. Consequently, the mediator and the IP of the interrupt controller (IC) were selected to be instantiated. As regards the memory management scheme, it was based on a *flat* address-space and therefore, no MMU components were instantiated in the system. The Figure 3(b) depicts the block diagram of the SOC that was generated after submit the application to the sequence of tools presented in section 5.

Figure 3. Block diagram of Leon2 (a) and the SoC that was experimentally customized (b).

Aiming at clarifying the expressiveness of this sample instance, it is important to compare the obtained results to the numbers of ordinary operating system that were ported to the LEON2 platform. Usually these systems are generated to compromise all the features that the SOC is able to provide. The absence of a component-based engineering and the lack of modern software engineering techniques affects not only the size and performance of the final system, but also effectively increase NRE costs and time-to-market. For instance, the EPOS instance generated to support the experimental application was 2.94 Kbytes and the associated SOC (Figure 3(b)) took 29% of the Virtex2 FPGA area. Traditional ports of UCLINUX [9] and ECOS [10] to the Leon2 platform have, respectively, 2,740 and 432.94 Kbytes each and the both were ported to a SoC instance with 13,261 LUTs (*Look-up-Table*), which represents 60% of the Virtex2 area.

6. CONCLUSION AND FUTURE WORK

In this article we conjectured about the the use of *hardware mediators* as a new co-design artifact. The deployment of AOSD in the context where mediators were originally proposed leaded us to use this portability artifact for the automatic generation of SoC-based embedded systems. The presented results are quite simple and just showed the viability of generating run-time support systems and system-on-chip instances considering the application's requirements. However, as exemplified in the section 4, we are not only able to identify which devices shall be instantiated in the SOC but, in fact, to configure each system component in order to better fits the application's requirements. In this sense a large gamma of new experiments started to be evaluated, such as processor scalability, memory hierarchy exploration and power management. The results obtained so far are encouraging and we hope present them soon.

References

[1] ARM (2003). *The de facto Standard for On-Chip Bus*. Advanced RISC Machines Limited, online document edition. http://www.arm.com/ products/ solutions/ AMBAHomePage.html.

[2] Bergamaschi, R. A., Bhattacharya, S., Wagner, R., Fellenz, C., Muhlada, M., Lee, W. R., White, F., and Daveau, J.-M. (2001). Automating the design of SoCs using cores. *IEEE Des. Test*, 18(5):32–45.

[3] Booch, G. (1994). *Object-Oriented Analysis and Design with Applications*. Addison-Wesley, 2nd edition.

[4] Frohlich, A. A. (2001). *Application-Oriented Operating Systems*. Number 17 in GMD Research Series. GMD - Forschungszentrum Informationstechnik, Sankt Augustin.

[5] Fröhlich, A. A. and Schröder-Preikschat, W. (1999). High performance application-oriented operating systems – the epos aproach. In *Proceedings of the 11th Symposium on Computer Architecture and High Performance Computing*, pages 3–9, Natal, Brazil.

[6] Fröhlich, A. A. and Schröder-Preikschat, W. (2000). Scenario adapters: Efficiently adapting components. In *Proceedings of the 4th World Multiconference on Systemics, Cybernetics and Informatics*, Orlando, U.S.A.

[7] Habermann, A. N., Flon, L., and Cooprider, L. (1976). Modularization and hierarchy in a family of operating systems. *Commun. ACM*, 19(5):266–272.

[8] Kiczales, G., Lamping, J., Mendhekar, A., Maeda, C., Lopes, C. V., Loingtier, J.-M., and Irwin, J. (1997). Aspect-oriented programming. In *Proceedings of the European Conference on Object-oriented Programming'97*, volume 1241 of *Lecture Notes in Computer Science*, pages 220–242, Jyv skyl, Finland. Springer.

[9] Wurm, M. (2003). uClinux for sparc-mmuless with ethernet mac. Technical report, Graz University of Technology.

[10] Massa, A. (2002). *Embedded SW. Development with eCos*. Prentice Hall, 1st edition.

[11] Mattsson, D. and Christensson, M. (2004). Evaluation of synthesizable cpu cores. Technical report, Chalmers University Of Technology.

[12] Mikhajlov, L. and Sekerinski, E. (1998). A study of the fragile base class problem. In *Proceedings of the 12th European Conference on Object-Oriented Programming*, pages 355–382. Springer-Verlag.

[13] Musser, D. R. and Stepanov, A. A. (1989). Generic programming. In *Proceedings of the First International Joint' Conference of ISSAC and AAECC*, number 358 in Lecture Notes in Computer Science, pages 13–25, Rome, Italy. Springer.

[14] Neville-Neil, G. and Whitney, T. (2003). Soc: Software, hardware, nightmare, bliss. *ACM Queue*, 1(2):24.

[15] Parnas, D. L. (1976). On the design and development of program families. *IEEE Transactions on Software Engineering*, SE-2(1):1–9.

[16] Polpeta, F. V. and Fröhlich, A. A. (2004). Hardware mediators: a portability artifact for component-based systems. In *Proceeding of the International Conference on Embedded and Ubiquitous Computing*, volume Lecture Notes in Computer Science, pages 271–280, Aizu,Japan. Springer.

[17] Rutenbar, R. A., Baron, M., Daniel, T., Jayaraman, R., Or-Bach, Z., Rose, J., and Sechen, C. (2001). (When) Will FPGAs kill ASICs? In *Proceedings of the 38th conference on Design automation*, pages 321–322. ACM Press.