

HYRA: A Software-defined Radio Architecture for Wireless Embedded Systems

Tiago Rogério Mück and Antônio Augusto Fröhlich
Federal University of Santa Catarina (UFSC)
Software/Hardware Integration Lab (LISHA)
Florianópolis - Brazil
{tiago, guto}@lisha.ufsc.br

Abstract—Traditional Software-defined Radio (SDR) architectures cannot go with the requirements of embedded systems, specially in terms of performance and power consumption. Low-power FPGAs now reaching the market might soon become a viable alternative to overcome such limitations. The *Hybrid Radio Architecture* (HYRA) introduced in this paper contributes to this scenario as it explores the *Hybrid HW/SW Component* concept to enable the implementation of SDRs as direct mappings of high-level synchronous data flow models. Although addressing SDR from a higher level of abstraction, HYRA mechanisms proved far more efficient than those behind GNU Radio when the target is an embedded reconfigurable hardware platform.

Keywords- *software-defined radio; embedded systems; FPGA.*

I. INTRODUCTION

Wireless communication devices are at the heart of a growing number of embedded systems. Some of them, such as smartphones, must implement multiple communication protocols in face of constantly evolving standards. Others, such as wireless sensor network gateways, must simultaneously communicate under multiple protocols and sometimes even dynamically adapt themselves to preserve connectivity. In this scenario, *Software-defined Radio* (SDR) becomes an appealing approach, since most of the key components in the communication system—including the physical layer—are pushed into software, thus making them easy to reconfigure [1].

However, the implementation of a wireless communication system based only on an RF front-end, A/D converters, and a *General-Purpose Processors* (GPP) comes at a high cost. The associated *Digital Signal Processing* (DSP) algorithms demand very high processing power, a requirement that contradicts major design premises in the field, which usually include low cost, low energy consumption, and small size. Nevertheless, it is important to notice that this exceeding demand for processing power arises basically from the serialization of essentially parallel algorithms that takes place as they are pushed from hardware to software.

Implementing the key concepts behind an SDR on a reconfigurable hardware platform such as an *Field Programmable Gate Array* (FPGA) would preserve its main advantage—flexibility—without requiring a high-performance processor. For instance, an architecture based on DSP blocks on a datapath implementing a *Synchronous Data Flow* (SDF) could take advantage of the platform's inherent parallelism for the

implementation of each individual DSP block and also to interconnect them efficiently. This has not been an option to embedded systems designers until now for a single reason: power consumption. Recent advances in low-power reconfigurable hardware, however, suggest that such systems can soon become viable. Indeed, several groups currently explore the use of hardware accelerators for the implementation of SDR algorithms [2]–[4]. *Single instruction, multiple data* (SIMD) extensions of GPPs, DSP processors, and functional blocks implemented in FPGAs are common approaches to limit processing power requirements at software level. Notwithstanding, the imminence of embedded SDRs fully implemented in low-power FPGAs calls for a systematic approach to guide the development of DSP components, interconnections, and controllers.

In this paper, we introduce HYRA, the *Hybrid Radio Architecture*, as a fundamental step toward a more comprehensive strategy to deploy SDRs in the context of embedded systems. HYRA relies on the *Hybrid HW/SW Component* concept of *Application-driven Embedded System Design* (ADESD) [5] to enable the implementation of an SDR as a direct mapping of a high-level SDF model. Each functional block in the model is associated to an hybrid component that can be plugged into HYRA's embedded SDR framework. Since hybrid components preserve their interfaces independently of how they are implemented, developers can freely decide which elements of the SDF graph go to software and which go to hardware. HYRA's framework features a programmable interconnect infrastructure that abstracts the *First In, First Out* (FIFO) channels between components. It also features a controller that dynamically coordinates the flow of data between components.

The remainder of this paper is organized as follows: Section II discusses related SDR implementation approaches; Section III recalls ADESD hybrid components, a fundamental concept behind HYRA; Section IV describes HYRA in details, while Section V presents an experimental evaluation; Section VI closes the paper with our conclusions.

II. RELATED WORK

The most straightforward SDR implementation approaches are the ones based on GPPs. These approaches target flexibility and ease of development, and usually delegate all processing to a GPP on a PC-like machine. These approaches

are not suitable to embedded systems not only because of cost, energy consumption, and size, but also because of the overhead imposed by general-purpose operating systems and by the high-latency, high-jitter communication interfaces used to reach the RF front-end [6]. The GNU Radio [7] is the most representative case in this group. It features a framework and a library of signal processing blocks that enables SDRs to be built on ordinary PCs. In GNU Radio, the physical layer of a radio is abstracted as a flow graph in which nodes represent processing blocks and edges represent the data flow between them.

Another approach is to delegate signal processing to programmable devices specifically designed for that purpose. Several architectures [2]–[4] relies on a GPP processor coordinating multiple DSP processors with SIMD and *very long instruction word* (VLIW) datapaths. Some architectures, such as the Elemental Computing Architecture [8], define fine-grained components specific to a given class of operations which can be configured and connected to each other to build an SDR. Differently from HYRA, these approaches focus on the efficient implementation of individual DSP blocks, without addressing the relationship between the implementation and a high level model. Also, the resource allocation and synchronization of processing elements must be controlled manually by the programmer. Some tools and languages, such as SPIR [9], aim to provide means to compile high level models of DSP applications into code that is suitable to run on *multiprocessor System-on-Chip* (MPSoC) DSP architectures. This is an important step toward a higher level SDR development strategy, but, as authors recognize, some algorithms that require considerably more processing power than the average, such as filters, searchers, and Turbo decoder, easily become bottlenecks in the programmable DSP hardware approach.

Apart from the software-based approaches, the dedicated implementation of the wireless communication protocols in FPGAs is another common approach to put together the flexibility of a reconfigurable radio, and the efficiency of a dedicated hardware. However, implementing a complex digital hardware design is not a straightforward task. Even if tools to translate high level specifications to a synthesizable *register transfer level* (RTL) description exist [10], [11], there is still a lack of means to integrate the dedicated hardware dataflow in a control flow that also encompasses software processes on the GPP. As a result, any change in the SDR protocol that requires more than a change on the parameters of existing hardware blocks usually requires the generation of a new hardware instance.

III. ADESD HYBRID COMPONENTS

HYRA was developed based on the idea of *hybrid hardware/software components* [5]. This concept is an elaboration on the concept of hardware mediators proposed in the *Application-driven Embedded System Design* (ADESD) [12] methodology. In ADESD, hardware mediators are a particular kind of component that are responsible for keeping the high-level abstractions independent of the hardware platform. This

components are implemented using *generative programming* techniques, adapting the hardware interface to the interface required by the system instead of creating a hardware abstraction layer. The idea of *hybrid hardware/software components* emerges from the fact that different mediators can exist for the same hardware component, each one designed for different purposes (e.g., changing the trade-off between performance and energy consumption). Each component aggregates mediators for its many implementations, which could be in hardware, software or both. According to system requirements of cost, performance, energy, etc., any one of these implementations can be selected without any change to the higher system layers that use the component.

In previous works [5] were defined and implemented in the *Embedded Parallel Operating System* (EPOS) [12] some architectural guidelines for the translation of operating system related components, such as timers, schedulers, and synchronizers, from software to hardware and vice versa. Whether such guidelines can also be defined for DSP related components has not yet been investigated, but nonetheless, a hybrid component is a convenient construct to encapsulate functional blocks on an SDR. This will be demonstrated in the next sections.

IV. SDR IMPLEMENTATION WITH HYRA

HYRA relies on SDF abstractions of SDRs. In this model, the SDR processing chain is abstracted as a flow graph, where the nodes represent processing blocks and the edges represent the data flow between the blocks. In HYRA, each functional block in the SDF is associated to a hybrid component that can be plugged into HYRA's embedded SDR framework. The developer uses this framework to specify connections that defines the data flow between components. The framework is responsible for creating the FIFO channels between the components and for starting the runtime mechanism that dynamically coordinates the data flow.

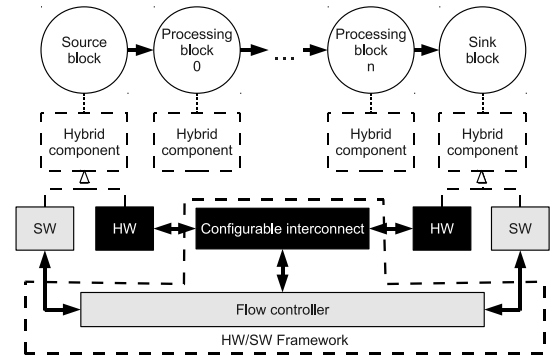


Fig. 1. Overview of HYRA

Figure 1 shows an overview of our architecture. Its framework has both a hardware and software side. The hardware side features a configurable interconnect structure which provides a FIFO-like stream interface to interconnect the hardware implementations of hybrid components. Also, it offers the

resources necessary to create HW FIFO channels between components in hardware and to coordinate their execution. The software side offers the interfaces to connect the software implementations of hybrid components and the runtime mechanism, denominated *Flow Controller*, which is responsible for controlling the connections and synchronization between components. The following sections will explain in more details each part of the architecture.

A. Flow controller

The *flow controller mechanism* is responsible for connecting the components and controlling the data flow between them at runtime. When the developer specifies a connection between two components, the flow controller creates a FIFO channel between them. The size of the FIFO in the channel is defined by the following equation:

$$FIFO_{size} = \max(Blk_0^{outputrate}, Blk_1^{inputrate}) \cdot \alpha \quad (1)$$

in which an output of Blk_0 is being connected to an input of Blk_1 , $Blk_0^{outputrate}$ is the number of data elements generated upon each execution of Blk_0 , and $Blk_1^{inputrate}$ is the number of data elements consumed in each execution of Blk_1 . The FIFO size is formulated in this way based on the fact that Blk_0 cannot generate data faster than Blk_1 can consume. If this happens in the system, due to modeling error or poor performance of Blk_1 , the FIFO will always overflow. α is a safety factor that should be set according to the jitter characteristics of the platform.

The FIFO allocation will depend on the actual physical implementation of the hybrid components that the channel is connecting. If both are implemented on software, a SW FIFO will be dynamically allocated in the system main memory. If one or both components are in hardware, the *flow controller mechanism* will allocate the FIFO inside the hardware interconnect structure. The next section will explain the hardware side of the framework.

The control of the data flow between software components is accomplished by creating a thread for each component. Each thread executes a loop where, at first, it remains locked onto semaphores associated with the channels connected to the block's inputs. Each time an element is added to a channel, the $v()$ method of its associated semaphore is called, unlocking the threads that consume the data from the channels. After acquiring all the semaphores, the thread consumes the inputs, executes the block's processing, and finally writes the result in the output channels, unlocking the threads associated to the subsequent blocks.

B. Hardware support

Hybrid components implemented in hardware don't use the software synchronization mechanism described in the previous section. Instead, they are controlled directly by signals provided by the FIFO channels in hardware. The deployment of HW FIFO channels is supported by the flow controller hardware structure shown in Figure 2. This structure mainly

consists of a interconnection block that have a set of read ports, write ports and internal FIFOs, where the connection between these three elements can be defined by software-controlled configuration registers.




Fig. 2. HW layer of HyRA's framework

All the components that are implemented as hardware have their inputs connected to the structure's read ports, and their outputs connected to the write ports. When these components are connected, the flow controller mechanism uses the information provided by the component's software interface to define which port must be connected to which FIFO. Since the HW FIFOs must have a fixed size, the interconnect structure allows FIFOs to be interconnected among them. This way, when two components are connected, it is possible to allocate a chain of FIFOs between them, in a way in which the total size of the chain is bigger than or equal to the required FIFO size.

Figure 3 shows how we have implemented this interconnect structure. We used a simplified butterfly fat tree NoC architecture [13] optimized for the interconnection of stream blocks. It consists of a matrix of FIFOs where each FIFO input is connected to each input port, and each output port is connected to each FIFO output. Each FIFO output is connected to the input of the FIFO in the next column on the same line. With this interconnect scheme we can provide a wide range of possible allocation for each input/output port, while keeping the use of FPGA resources by interconnect at a reasonable level.

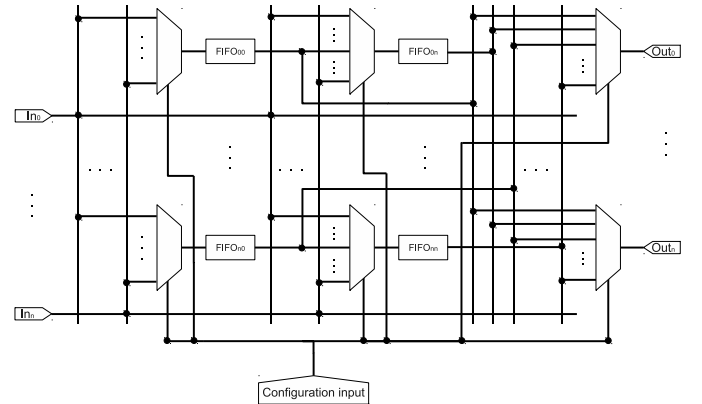


Fig. 3. Overview of the flow controller HW interconnect internal structure

In order to provide all possible connections between HW

and SW components, four different FIFO channel implementation are required: *SW FIFO*, *HW FIFO*, *SW-HW FIFO*, and *HW-SW FIFO*. We already explained how channels between SW-only and HW-only components are created. The connection between software and hardware components is achieved through the *Block SW Interface* shown in Figure 2. It behaves like a wrapper between the system bus and the interconnect structure stream interface. When a component in hardware is connected to a component in software, its respective ports are connected to ports associated with the *Block SW Interface*. When there is a SW→HW connection a *SW-HW FIFO* provides a software interface so the source block can write to the *Block SW Interface* output port associated to the destination block input port. But, the HW→SW connection requires additional runtime and hardware support. When a *HW-SW FIFO* is created, it register itself in the interrupt handler for the *Block SW Interface*'s interrupts. Every time new data arrives at one of the FIFOs connected to the *Block SW Interface*'s input ports, it will issue an interrupt that will release the semaphore associated with the FIFO, as described in the previous section.

V. EVALUATION AND RESULTS

In this work, we have focused only on the architectural support and data flow aspects for the implementation of SDR. In order to evaluate the proposed architecture we disconsidered the signal processing function, since they are covered in other works [10], [11], and focus only on HYRA's intrinsic overhead. We have evaluated this overhead in two aspects: area overhead (FPGA resource utilization) and performance overhead (latency added to the data flow by the hardware and software control structures).

A. Evaluation setup

To define the data flow structures for our evaluations, we have analyzed the data flow structure of the physical layer of several protocols covering a wide range of modulation schemes and application classes: Bluetooth, UWB, ZigBee, Wi-Fi (802.11a), and W-CDMA. In this analysis we verified that all protocols follows a common data flow structure. On the receive chain, there is usually a filter before the demodulation blocks, normally a low pass filter used to obtain a clean piece of spectrum that contains the information. Next, there are the demodulation/synchronization blocks, which normally consists of one or more data flows being processed in parallel. The last step is a post-demodulation filter which normally consists of a channel decoder for error detection and correction. The transmit chain follows an analogous structure. From this general structure, we have defined the structures shown in Figure 4 to evaluate the overhead in terms of the number of blocks in a data flow (4a) and the number of data flows in parallel (4b), covering many possible variations of the general structure described previously. We also analyzed how the number of inputs/outputs of a block affects the overhead (4c).

These structures are composed by three kinds of blocks. The *Timestamp source* block generates samples which consist

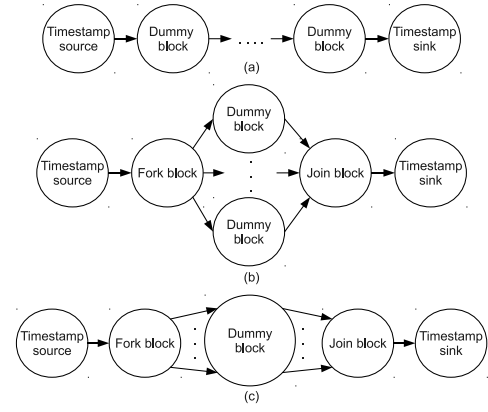


Fig. 4. SDRs data flow structures used for the overhead evaluation in terms of the number of blocks in serial (a), the number of data flows in parallel (b), and the number of inputs/outputs (c)

of timestamps that represent the time when the sample was generated. The *dummy blocks* are empty blocks that just propagate their inputs to their outputs. After being generated, the samples will go through the *dummy block* chain. When a sample arrives at the *Timestamp sink* block, the timestamp is compared with the current time, obtaining the time the sample took to go through the *dummy block* chain. Since the *dummy blocks* are empty, this resulting time represents only the overhead imposed by the architecture on the data flow. There is also the *Fork block* and *Join block* which are used to fork and join the data flows, respectively.

B. System implementation and configuration

To evaluate these three structures, we have implemented HYRA on the EPOS operating system running on the Xilinx's ML403 Embedded Platform. The ML403 features a Virtex-4 FPGA with an embedded PowerPC 405 microprocessor. In order to use the same hardware configuration, we have synthesized the hardware with all of the necessary blocks for all experiments. We have used the following tools and parameters: ISE/EDK 10.1; GCC 4.0.2; FPGA and microprocessor clocked at 100 MHz; interconnect structure configured with 32 input/output ports, and 64 FIFOs (8 bit wide with 16 elements); the α factor fixed to 1 in order to provide an evaluation considering low jitter requirements.

Table I shows the resource consumption of the generated hardware. Separate results are shown for HYRA's structures along with the HW dummy blocks, and for the system IPs generated by EDK (internal memory, memory controller, interruption controller, UART, etc). Our architecture alone uses about 65% of the available logic and 0% of the available memory. Due to the lack of memory blocks available on the device, we chose to implement the FIFOs using the SRL16 capabilities to convert a 4-input LUT into a 16-bit shift register. However, this apparently high resource usage is due to the very limited amount of logic available on the used device. When compared to other system IPs, we can see that HYRA uses slightly more resources than a complete set of basic IO and memory peripherals.

TABLE I
AMOUNT OF HW RESOURCES USED BY THE SYNTHESIZED STRUCTURES

Resource	Our IPs	EDK IPs	Full System
4-input LUTs	37%	35%	72%
Slice Flip Flops	67%	31%	98%
Occupied Slices	70%	55%	99%
RAM blocks	0%	63%	63%
Max. frequency	167 MHz	109 MHz	107 MHz

C. Performance overhead

We have done tests to determine the performance overhead that we have defined as the latency between the *Timestamp source* and *sink* blocks in the three basic data flow structures. We have implemented each structure in hardware and software and executed tests with the number of dummy blocks ranging from 1 to 16. In each test 6×10^7 samples were generated and we obtained the average value of the latencies of each sample and the standard deviation which was used to obtain the coefficient of variation. A sampling rate of 1×10^6 samples/second was used in the tests with blocks in hardware. For the tests with blocks in software we used a sampling rate of 1×10^4 due to the low speed of the PowerPC processor.

Figures 5a and 5b show the results. When using only software blocks, the overhead grows linearly in relation to the increase in the number of blocks and the number of inputs and outputs for all structures. The coefficient of variation remained low in all configurations. When using only HW blocks, the latency was about four orders of magnitude lower than when using software blocks and, as expected, except for the serial block configuration, the latency remained constant regardless the size of the structure, due to the full parallelism that can be explored in this kind of architecture. There is also a null coefficient of variation in the hardware operations.

To evaluate the communication latency between components implemented in HW and components implemented in SW, we have used the data flow shown in Figure 6 which cover operations on both *SW-HW FIFOs* and *HW-SW FIFOs*. We have performed the same experiment described previously on this two structures and verified that the average latency on both interleaved data flows yielded similar results: $221\mu s$ and $234\mu s$ for data flows (a) and (b), respectively. Data flow (a) have more SW blocks then (b), thus showing a higher SW management overhead. However, both have two *SW-HW FIFOs* and two *HW-SW FIFOs* connecting the blocks. By comparing the latencies with the ones obtained for SW-only blocks ($221\mu s$) and HW-only blocks ($0.31\mu s$) we can see that the read/write operations SW channels represents the most significant overhead.

We also compared the overhead of our architecture to GNU Radio. For this comparison, we replicated the same tests described previously using GNU Radio running over a Linux operating system in a PC. Our architecture and EPOS were compiled for the IA32 architecture only with software blocks support, and evaluated in the same system. For the GNU

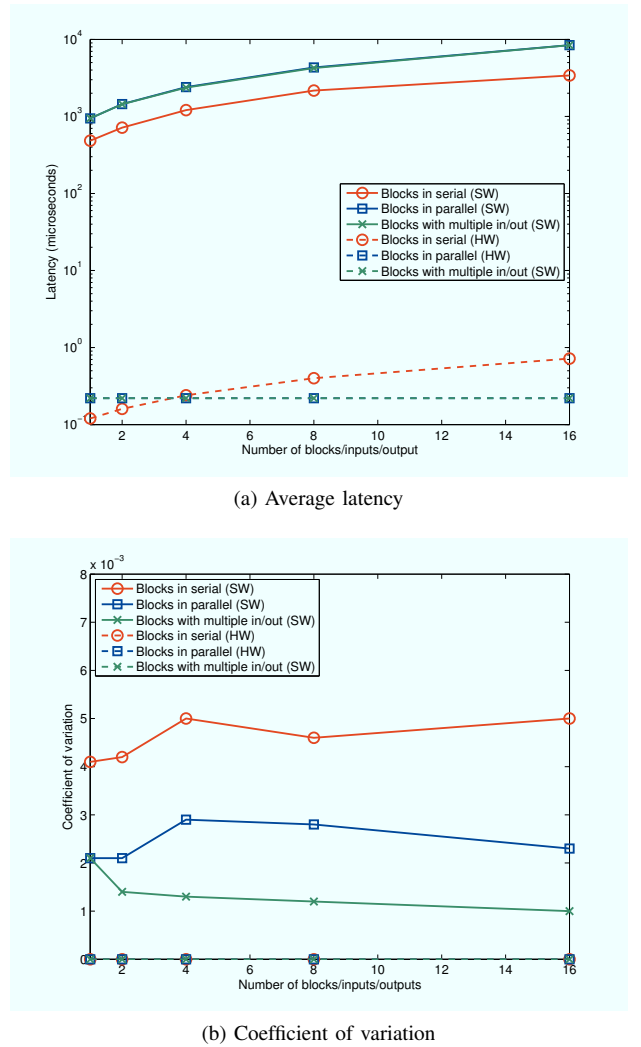


Fig. 5. Latency for blocks in serial, in parallel, and with multiple input/output





Fig. 6. Serial data flows with interleaved SW and HW blocks


Radio experiment, we used GNU Radio 3.2.2 running on a Linux kernel 2.6.28. The result for the serial blocks data flow structure shown in Figure 7a demonstrates that our architecture performance surpasses GNU Radio between 2 and 4 times, and this difference increases as the number of blocks in the processing chain increases. Figure 7b also shows that we are able to achieve smaller latency variations as well.

D. Discussion

The results show that our architecture yields superior performance than an equivalent, in terms of abstraction level,



(a) Average latency



(b) Coefficient of variation

Fig. 7. Latency of the proposed architecture VS GNU Radio on the serial blocks data flow structure

commonly used architecture. Only software components were used in this comparison, if a hardware device was used to generate the timestamps, GNU Radio would suffer an additional disadvantage. In GNU Radio, the use of any hardware device to obtain or sink data from/to the environment requires Linux drivers whose performance is mostly limited by the kernel's abstraction layer. A previous work [6] shows that, due to the Linux kernel overhead, the standard deviation of the time a sample takes to get to the processing chain after being generated in the RF Front-end is higher than the average time. This problem does not appear in EPOS since the metaprogrammed hardware mediators are dissolved within the application when the system is compiled, which leads to higher performance.

However, even with known latency problem, the GNU Radio is widely used and several protocols have been successfully implemented using it. The results have shown that with our architecture we were able to bring similar functionality with superior performance to the embedded system domain, which

leads to the conclusion that our architecture is suitable for the implementation of high-end protocols in embedded systems.

VI. CONCLUSION

In this paper we have introduced HYRA, an *Hybrid Radio Architecture* that explores the *Hybrid Component* concept within ADESD to enable the implementation of SDRs as direct mappings of high-level SDF models. As hybrid components, HYRA SDR blocks can be implemented as arbitrary combination of software and hardware on FPGA-based platforms. The programmable interconnect infrastructure in HYRA's framework ensures transparency in this respect. FIFO channels can be fine tuned to fulfill the requirements of a given SDR protocol, while the controller dynamically coordinates the flow of data between components.

In comparison with other approaches, HYRA addresses the implementation of SDRs in the context of embedded systems from a higher level of abstraction. Moreover, the evaluation results presented in this paper confirm that the overhead caused by the proposed architecture in terms of latency is much smaller than that of GNU Radio, a widely accepted architecture. Furthermore, our experiments demonstrated that HYRA can be implemented on reconfigurable hardware platform with minimal additional resources. In combination, this factors confirm that our architecture meet the requirements for the implementation of high-end protocols in embedded systems.

REFERENCES

- [1] E. Buracchini, "The software radio concept," *IEEE Comm. Mag.*, vol. 38, no. 9, pp. 138–143, 2000.
- [2] J. Glossner, E. Hokenek, and M. Moudgill, "The Sandbridge Sandblaster Communications Processor," in *3rd WASP*, 2004, pp. 53–58.
- [3] M. Woh, Y. Lin, S. Seo, S. Mahlke, T. Mudge, C. Chakrabarti, R. Bruce, D. Kershaw, A. Reid, M. Wilder, and K. Flautner, "From SODA to scotch: The evolution of a wireless baseband processor," in *MICRO 41*. Washington, DC, USA: IEEE Computer Society, 2008, pp. 152–163.
- [4] K. van Berkel, F. Heinle, P. P. E. Meuwissen, K. Moerman, and M. Weiss, "Vector processing as an enabler for software-defined radio in handheld devices," *EURASIP*, vol. 2005, pp. 2613–2625, 2005.
- [5] H. Marcondes and A. A. Fröhlich, "A Hybrid Hardware and Software Component Architecture for Embedded System Design," in *IESS '09*, Langenargen, Germany, 2009, pp. 259–270.
- [6] G. Nychis, T. Hottelier, Z. Yang, S. Seshan, and P. Steenkiste, "Enabling MAC Protocols Implementation on Software-defined Radios," in *NSDI '09*, 2009.
- [7] GNU FSF project, "The GNU Radio," 2010, [retrieved, November 8, 2010]. [Online]. Available: <http://www.gnu.org/software/gnuradio>
- [8] Steven Kelem, Brian Box, Stephen Wasson, Robert Plunkett, Joseph Hassoun, and Chris Phillips, "An Elemental Computing Architecture for SD Radio," in *SDR '07*, 2007.
- [9] Y. Lin, M. Kudlur, S. Mahlke, and T. Mudge, "Hierarchical coarse-grained stream compilation for software defined radio," in *CASES '07*. New York, NY, USA: ACM, 2007, pp. 115–124.
- [10] Z. Guo, W. Najjar, and B. Buyukkurt, "Efficient hardware code generation for FPGAs," *ACM TACO*, vol. 5, no. 1, pp. 1–26, 2008.
- [11] F. Plavec, Z. Vranesic, and S. Brown, "Towards compilation of streaming programs into FPGA hardware," in *FDL '08*, 2008, pp. 67–72.
- [12] A. A. Fröhlich, *Application-Oriented Operating Systems*, ser. GMD Research Series. Sankt Augustin: GMD - Forschungszentrum Informationstechnik, Aug. 2001, no. 17.
- [13] P. P. Pande, C. Grecu, A. Ivanov, and R. Saleh, "Design of a switch for network on chip applications," in *ISCAS '03*, vol. 5, 2003, pp. V217–V220.