

A Software-defined Radio Architecture for Embedded Systems

Abstract

Traditional Software-defined Radio architectures cannot go with the requirements of embedded systems, specially in terms of performance and power consumption. Low-power FPGAs now reaching the market might soon become a viable alternative to overcome such limitations. The *Hybrid Radio Architecture* (HYRA) introduced in this paper contributes to this scenario as it explores the *Hybrid HW/SW Component* concept to enable the implementation of SDRs as direct mappings of high-level SDF models. Although addressing SDR from a higher level of abstraction, HYRA mechanisms proved far more efficient than those behind GNU Radio when the target is a embedded reconfigurable hardware platform.

1 Introduction

Wireless communication devices are at the heart of a growing number of embedded systems. Some of them, such as smartphones, must implement multiple communication protocols (e.g. GSM, GPRS, UMTS, Wi-Fi, Bluetooth) in face of constantly evolving standards. Others, such as wireless sensor network gateways, must simultaneously communicate under multiple protocols and sometimes even dynamically adapt themselves to preserve connectivity. In this scenario, *Software-defined Radio* (SDR) becomes an appealing approach, since most of the key components in the communication system—including the physical layer—are pushed into software, thus making them easy to reconfigure [1].

However, the implementation of a wireless communication system based only on an RF front-end, A/D converters, and a processor as illustrated in figure 1 comes at a high cost. The associated *Digital Signal Processing* (DSP) algorithms demand very high processing power, a requirement that contradicts major design premises in the field, which usually include low cost, low energy consumption, and small size. Nevertheless, it is important to notice that this exceeding demand for processing power arises basically from the serialization of essentially parallel algorithms that takes place as they are pushed from hardware to software, and from the high-latency datapath typical of software-oriented architectures.

Implementing the key concepts behind an SDR on a reconfigurable hardware platform such as an FPGA would preserve its main advantage—flexibility—without requiring a high-performance processor. For in-




Figure 1: General architecture of an SDR.

stance, an architecture based on DSP blocks on a datapath implementing a *Synchronous Data Flow* (SDF) [2] could take advantage of the platform’s inherent parallelism for the implementation of each individual DSP block and also to interconnect them efficiently. This has not been an option to embedded systems designers until now for a single reason: power consumption. Recent advances in low-power reconfigurable hardware, however, suggest that such systems can soon become viable. Indeed, several groups currently explore the use of hardware accelerators for the implementation of SDR algorithms [3,4]. SIMD extensions of general-purpose processors, DSP processors, and functional blocks implemented in FPGAs are common approaches to limit processing power requirements at software level. Notwithstanding, the imminence of embedded SDRs fully implemented in low-power FPGAs calls for a systematic approach to guide the development of DSP components, interconnections, and controllers.

In this paper we introduce HYRA, the *Hybrid Radio Architecture*, as a fundamental step toward a more comprehensive strategy to deploy SDRs in the context of embedded systems. HYRA relies on the *Hybrid HW/SW Component* concept of *Application-driven Embedded System Design* (ADESD) [5] to enable the implementation of an SDR as a direct mapping of a high-level SDF model. Each functional block in the model is associated to an hybrid component that can be plugged into HYRA’s embedded SDR framework. Since hybrid components preserve their interfaces independently of how they are implemented (hardware-only, software-only, or a hardware/software mix), developers can freely decide which elements of the SDF graph go to software and which go to hardware. HYRA’s framework features a programmable interconnect infrastructure that abstracts the FIFO channels between components. It also features a controller that dynamically coordinates the flow of data between components.

The remainder of this paper is organized as follows: section 2 discusses related SDR implementation approaches; Section 3 recalls ADESD hybrid components, a fundamental concept behind HYRA; Section 4 describes HYRA in details, while 5 presents an experimental evaluation; Section 6 closes the paper with our conclusions.

2 Related Work

SDR approaches based on *General-Purpose Processors* (GPP) target flexibility and ease of development. They usually strongly adhere SDR theoretical principles, delegating all processing to the GPP on a PC-like machine. These approaches are not suitable to embedded systems not only because of cost, energy consumption, and size, but also because of the overhead imposed by general-purpose operating systems and by the high-latency, high-jitter communication interfaces used to reach the RF front-end [6]. The GNU Radio [7] is the most representative case in this group. It features a framework and a library of signal processing blocks that enables SDRs to be built on ordinary PCs easily and quickly. In GNU Radio, the physical layer of a radio is abstracted as a flow graph in which nodes represent processing blocks and edges represent the data flow between them.

Another approach is to delegate signal processing to programmable devices specifically designed for that purpose, including DSPs, GPPs with SIMD extensions, and fully dedicated devices such as digital up/down converters (DUC/DDC). The usage of FPGAs to implement high data rate functions (e.g. decimation, interpolation, and translation) that are integral to the processing chain of many protocols is also a trend in this scenario. It is well represented by the Universal Software Radio Peripheral (USRP) [8].

The Sandblaster architecture [9] relies on multiple DSP processors on a SIMD datapath to implement a sort of vector processing engine. A customized C compiler automatically maps vector operations to threads on the DSPs, enabling efficient architectural exploration without the need of low-level programming. A companion ARM core takes care of general-purpose tasks. The architecture's ability to handle heavy protocols was confirmed by a full implementation of W-CDMA [10]. Differently from the architecture proposed here, Sandblaster focus on the efficient implementation of individual DSP blocks, without addressing the relationship between them and an SDF model.

SODA defines an architecture that explores SIMD parallelism on a hardware optimized for 16 and 8 bits computations. It consists of a set of processing elements, each featuring both scalar and vectorial units with dedicated scratch-pad memories for instructions and data. An ARM core is used to implement higher protocol layers and also to synchronize the computations on each PE [11, 12]. An instance of SODA with four processing elements at 65nm CMOS was reported to meet the throughput requirements of W-CDMA and 802.11a protocols under an energy budget that makes it suitable to many embedded systems. However, SODA's processing elements must be programmed in assembly and the allocation of scratch-pad memories

must be controlled manually by the programmer. Support for coordinating the PEs based on an abstract functional model was also not yet addressed.

EVP is a *Very Long Instruction Word* (VLIW) architecture that features both a scalar and a SIMD datapath [13]. EVP-C, an extension of C used in the project, allows programmers to specify the mapping of software functional units into the architecture's elements. It also features vector data types and functions, performing automatic register allocation as well as VLIW instruction scheduling. Despite the use of a dedicated programming environment, EVP does not support higher level abstract models and published figures reveals limitations to support high data rate algorithms (e.g. Turbo decoding) that are delegated to additional hardware accelerators.

The Elemental Computing Architecture [14] defines “Elements” as fine-grained components specific to a given class of operations, such as ALU, barrel shifter, and state machine engine. An element can have several active execution contexts, what makes it a multitasking device. Elements are connected hierarchically, thus giving scalability to the architecture. Local zones are shaped by directly connecting elements that are subsequently interconnected through queues. The architecture performance has been demonstrated by an AES cipher and a 4096-point FFT, suggesting it is able to support high-end protocols. However, the implementation of elements and their coordination is not directly addressed by the architecture.

SPIR is a data flow language for SDR development [15]. It includes a set of tools that is able to translate a SPIR specification into code that is suitable to run on MPSoC DSP architectures such as SODA and Sandblaster. SPIR handles the scheduling and synchronization of processing blocks that have been previously implemented for the target DSP architecture. It follows a software pipelining technique inspired on *modulo scheduling* and uses *Integer Linear Programming* (ILP) to optimize the allocation of blocks in terms of memory and time [16]. SPIR is an important step toward a higher level SDR development strategy, but, as authors recognize, algorithms that require considerably more processing power than the average, such as filters, searchers, and Turbo decoder, easily become bottlenecks in the architecture.

In general, the approaches combining GPP and dedicated hardware yield a good trade-off between performance and the traditional embedded system design goals. However, since they rely of fixed hardware units, the partitioning of memory intense process across parallel units often becomes a pitfall. Not rarely, this must be done by hand, along with the scheduling and synchronization of functional blocks. Even if tools to translate high level specifications to a synthesizable RTL description exist [17–20], there is still a lack of means to integrate the dedicated hardware dataflow in a control flow that also encompasses software

processes on the GPP. As a result, any change in the SDR protocol that requires more than a change on the parameters of existing hardware blocks usually requires the generation of a new hardware instance.

3 ADESD Hybrid Components

Application-driven Embedded System Design (ADESD) elaborates on commonality and variability analysis—the well-known domain decomposition strategy behind Object-Orientation—to add the concept of aspect identification and separation at early stages of design [21]. It defines a domain engineering strategy focused on the production of families of scenario-independent components. Dependencies observed during domain engineering are captured as separate *aspect programs*, thus enabling components to be reused on a variety of execution scenarios with the application of proper aspect programs.

In this context, *Hardware Mediators* are a key concept in ADESD to ensure component portability across distinct hardware platforms. They define strict hardware/software interfaces for hardware components at a level that makes it possible to encompass quite different devices under a common interface. Devices missing features are complemented by software in a way similar to *Hardware Abstraction Layers* (HAL) in ordinary operating systems. However, differently from a HAL, a set of hardware mediators do not define a layer, since they are meant to be implemented as metaprograms (e.g. C++ templates with inline assembly) and thus get dissolved within higher-level components as soon as the interface contract is met. In this way, mediators can avoid a large fraction of the overhead associated by HALs, particularly in respect unneeded functionality and function call indirection. Additionally, hardware mediators constitute an efficient way to encapsulate synthesizable hardware components in programmable platforms, since they can be adjusted to finely match the corresponding glue logic.

The realization that multiple hardware mediators can coexist in the repository to represent distinct mixes of software and hardware used in the implementation of multiple versions of a given synthesizable component led ADESD to the concept of *Hybrid Components*. Indeed, ADESD even defines architectural guidelines for the translation of operating system related components, such as timers, schedulers, and synchronizers, from software to hardware and vice versa [5]. Whether such guidelines can also be defined for DSP related components has not yet been investigated, but nonetheless, a hybrid component is a convenient construct to encapsulate functional blocks on an SDR. This will be demonstrated in the next sections.

4 SDR Implementation with HYRA

HYRA relies on SDF abstractions of SDRs. In this model, the SDR processing chain is abstracted as a flow graph, where the nodes represent processing blocks and the edges represent the data flow between the blocks. In HYRA, each functional block in the SDF is associated to a hybrid component that can be plugged into HYRA's embedded SDR framework. The developer uses this framework to specify the connections that defines the data flow between the components. The framework is responsible for creating the FIFO channels between the components and for starting the run time mechanism that dynamically coordinates the data flow between components.

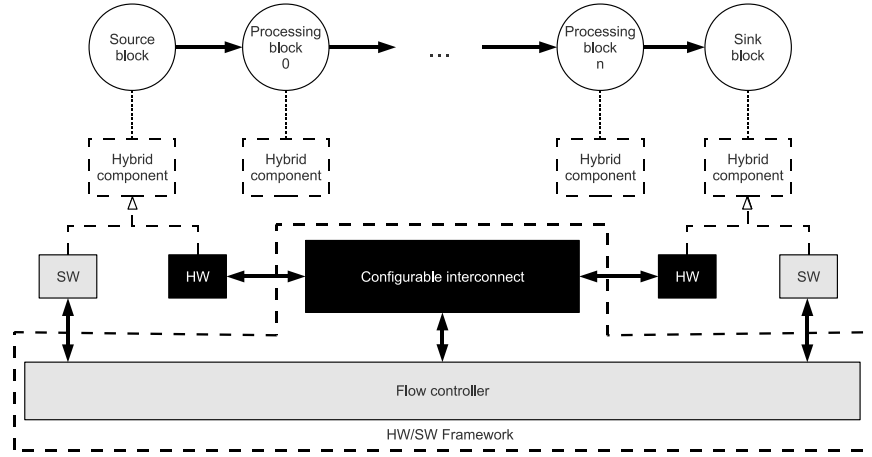


Figure 2: Overview of the proposed HYRA

Figure 2 shows an overview of our architecture. It's framework have both a hardware and software side. The hardware side features a configurable interconnect structure that provides a FIFO-like stream interface to connect the hardware implementations of hybrid components. Also, it offers the resources necessary to create HW FIFO channels between components in hardware and to coordinate their execution . The software side offers the interfaces to connect software implementations of hybrid components and the run-time mechanism, denominated *Flow Controller*, which is responsible for controlling the connections and synchronization between the components. The following sections will explain in more details each part of the architecture.

4.1 Flow Controller

The *flow controller mechanism* is responsible for connecting the components and controlling the data flow between them at run time. When the developer specifies a connection between two components, it first checks the component's stream interfaces to determine if one component generates the same data type that the other is expecting to consume. Then, it creates a FIFO channel to connect the components. The size of the FIFO in the channel is defined by the following equation:

$$FIFO_{size} = \max(Blk_0^{outputrate}, Blk_1^{inputrate}) \cdot \alpha \quad (1)$$

Where an output of Blk_0 is being connected to an input of Blk_1 , $Blk_0^{outputrate}$ is the number of data elements generated upon each execution of Blk_0 , and $Blk_1^{inputrate}$ is the number of data elements consumed in each execution of Blk_1 . The FIFO size is formulated in this way based on the fact that Blk_0 cannot generate data faster than Blk_1 can consume. If this happens in the system, due to modeling error or poor performance of Blk_1 , the FIFO will always overflow. α is a safety factor that should be set according to the jitter characteristics of the platform, in order to make sure the FIFO will be big enough to handle all the data generated by Blk_0 until it can be consumed by Blk_1 .

The FIFO allocation will depend on the actual physical implementation of the hybrid components that the channel is connecting. If both are implemented on software, the FIFO will be dynamically allocated in the system main memory. If one or both components are in hardware, the *flow controller mechanism* will allocate the FIFO inside the hardware interconnect structure. The next section will explain the hardware side of the framework.

The control of the data flow between software components is accomplished by creating a thread for each component, where its function is executed. The threads are synchronized using semaphores associated with the FIFO channels. At first, the threads remain locked onto the semaphores associated with the channels connected to the block's inputs. Each time an element is added to a channel, the $v()$ method of its associated semaphore is called, unlocking the threads that consume the data from the channels. The pseudo-code in figure 3 shows this behavior for a component associated to processing block. $In_{0..n}$ and $Out_{0..n}$ are the channels associated with the component's inputs and outputs, respectively. The semaphore is encapsulated in the FIFO channel and the *wait* and *signal* methods are related to the p and v operations of the semaphore.

```

Processing block loop :
.   In0.wait()
.   ...
.   Inn.wait()
.   If there are enough inputs:
.       .   Consume inputs
.       .   Do processing
.       .   For each output generated:
.       .       .   Write outputs to Out0
.       .       .   ...
.       .       .   Write outputs to Outn
.       .       .   Out0.signal()
.       .       .   ....
.       .       .   Outn.signal()

```

Figure 3: Data flow control of a software component representing a processing block

4.2 Hardware Support

Hybrid components implemented in hardware don't use the software synchronization mechanism described in the previous section. Instead, they are controlled directly by signals provided by the FIFO channels in hardware. The deployment of HW FIFO channels is supported by the flow controller hardware structure shown in figure 4. This structure mainly consists of a interconnection block that have a set of read ports, write ports and internal FIFOs, where the connection between these three elements can be defined by software-controlled configuration registers.




Figure 4: HW layer of HYRA's framework

All the components that are implemented as hardware have their inputs connected to the structure's read ports, and their outputs connected to the write ports. When these components are connected, the flow controller mechanism uses the information provided by the component's software interface to define which port must be connected to which FIFO. Since the HW FIFOs must have a fixed size, the interconnect

structure allows FIFOs to be interconnected among them. This way, when two components are connected, it is possible to allocate a chain of FIFOs between them, in a way in which the total size of the chain is bigger than or equal to the required FIFO size, as defined in section 4.1.

Figure 5 shows how we have implemented this interconnect structure. We used a simplified butterfly fat tree NoC architecture [22] optimized for the interconnection of stream blocks. It consists of a matrix of FIFOs where each FIFO input is connected to each input port, and each output port is connected to each FIFO output. Each FIFO output is connected to the input of the FIFO in the next column on the same line. With this interconnect scheme we can provide a wide range of possible allocation for each input/output port, while keeping the use of FPGA resources by interconnect at a reasonable level. This interconnect structure was designed to be fully parameterizable. We are able to choose the size of the FIFO matrix in both dimensions, as well as the FIFOs width and depth, so the synthesized hardware can be adapted according to the target platform.

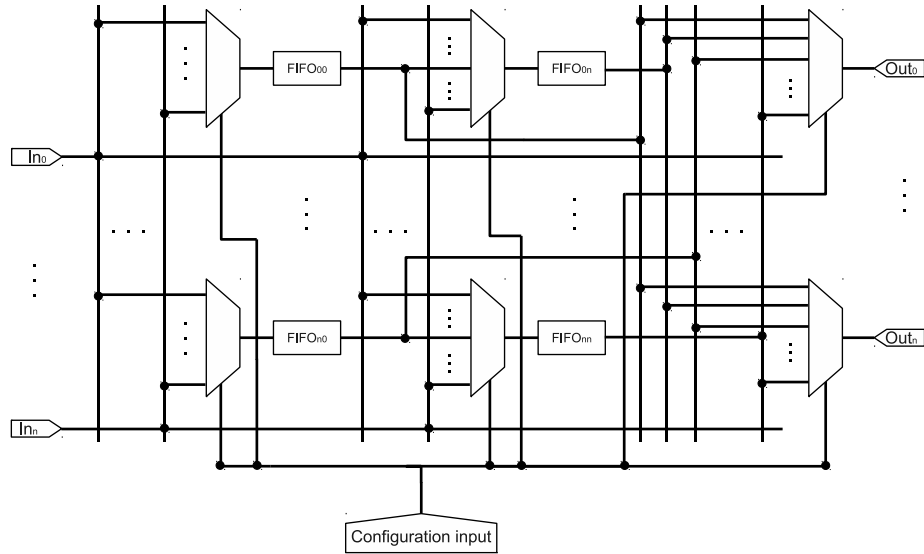


Figure 5: Overview of the flow controller HW interconnect internal structure

In order to provide all possible connections between HW and SW components, four different FIFO channel implementation are required: *SW FIFO*, *HW FIFO*, *SW-HW FIFO*, and *HW-SW FIFO*. We already explained how channels between components in the same domain are created. The connection between software and hardware components is achieved through the *Block SW Interface* shown in figure 4. It behaves like a wrapper between the system bus and the interconnect structure stream interface. When a component in hardware is connected to a component in software, its respective ports are connected to ports associated

with the *Block SW Interface*.

When there is a SW→HW connection a *SW-HW FIFO* provides a software interface so the source block can write to the *Block SW Interface* output port associated to the destination block input port. The HW→SW connection requires additional runtime and hardware support. When a *HW-SW FIFO* is created, it register itself in the interrupt handler for the *Block SW Interface*'s interrupts. Every time new data arrives at one of the FIFOs connected to the *Block SW Interface*'s input ports, it will issue an interrupt that will release the semaphore associated with the FIFO, as described in the previous section.

4.3 Deployment Example

This section presents a simple example that employs HYRA. Figure 6 shows the SDF of a simple spectrum analyser and the code that implements it in our architecture. The source block is a RF front-end that samples some window of the spectrum. The samples go through a DDC, which makes frequency translation and decimation so as to select the channel with the specified center frequency and bandwidth. An FFT convert the signal to the frequency domain and sends the information to a display sink that will plot the data on the screen. We consider that the FFT has an implementation in both hardware and software. The implementation that will be used is selected using *configurable traits*, which consists of compile-time data structures that contains information about objects in the system.

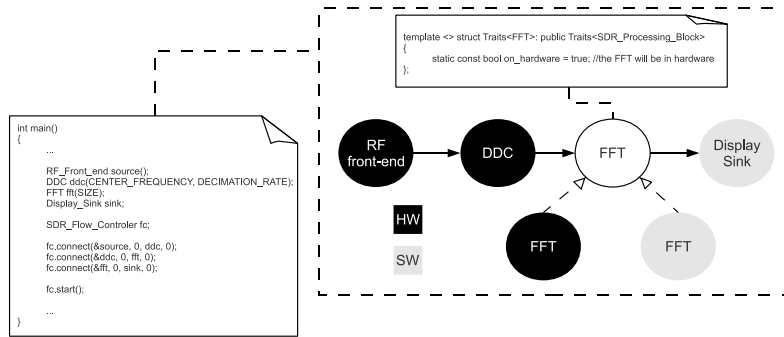


Figure 6: Simple spectrum analyser implemented using our architecture

Considering the scenario defined by the code in figure 6. When the RF front-end is connected to the DDC, the flow controller will configure the hardware interconnect structure to connect the ports associated to the blocks output and input, respectively, to a free FIFO, creating a *HW FIFO* channel between them. The same happens between the DDC and the FFT, which is in hardware. The information about the ports in which the components are connected in hardware are also defined in its respective *configurable traits*. Now,

when the FFT is connected to the display sink, which is a software block, the flow controller will follow the procedure described in the previous section, allocating a free port in the *Block SW Interface* and connecting it to the FIFO connected to the FFT's output port. If the application is compiled with the option *on_hardware = false* in the FFT's traits, the FFT will be implemented in software and a *HW-SW FIFO* will be created between the DDC and the FFT. Now, a simple *SW FIFO* is allocated between the FFT and the display sink.

5 Evaluation and results

In this work we have focused only on the architectural support and data flow aspects for the implementation of SDR. So, in order to evaluate the proposed architecture we disconsidered the signal processing function, since they are covered in other works [17, 18, 20, 23], and focus only on HYRA's intrinsic overhead. This overhead can be evaluated in two aspects: area overhead (FPGA resource utilization and code size) and performance overhead (latency added to the data flow by the hardware and software control structures). In the next sections we provide an analysis of data flow structures for various protocol categories, which allowed us to extract the basic types of data flow structures used in the subsequent evaluations.

5.1 Evaluation setup

So as to define the data flow structures for our evaluations, we have analyzed the SDR implementation and the data flow structure of the physical layer of several protocols: IEEE 802.15.1 (Bluetooth) [24], IEEE 802.15.3 (UWB) [25], IEEE 802.15.4 (ZigBee) [26], IEEE 802.11a (Wi-Fi) [11], IEEE 802.16 [27], and W-CDMA (UMTS) [11, 28]. With this protocols we can cover a wide range of modulation schemes and very distinct application classes: low data rate WPANs (IEEE 802.15.1 and IEEE 802.15.4), high data rate WPANs (IEEE 802.15.3), WLANs (IEEE 802.11a), and long range networks like WMANs (IEEE 802.16) and cellphone networks (W-CDMA). In this analysis we verified that the protocols generally follow the structure in figure 7. On the receive chain, there is usually a filter before the demodulation blocks, normally a low pass filter used to obtain a clean piece of spectrum that contains the information. Next are the demodulation/synchronization blocks, which normally consists of one or more data flows being processed in parallel. The last step is a post-demodulation filter which normally consists of a channel decoder for error detection and correction. The transmit chain follows an analogous structure.

From the general structure in figure 7 we have defined the structures shown in figure 8 to evaluate the




Figure 7: Common data flow structure for the receive chain of most wireless protocols

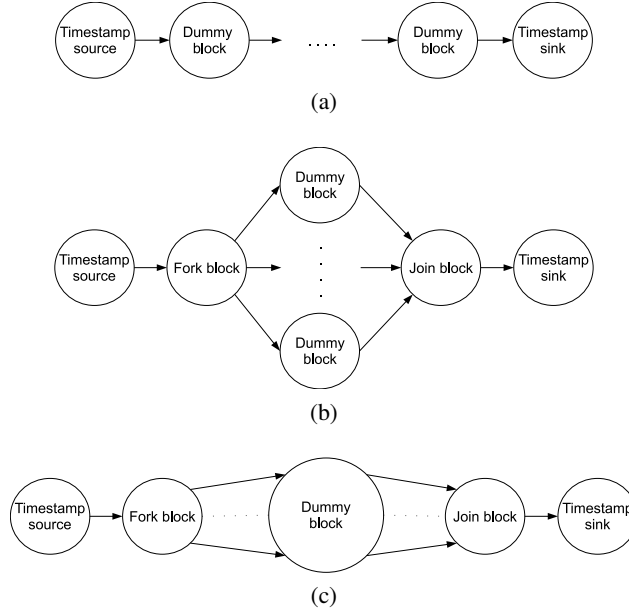


Figure 8: SDRs data flow structures used for the overhead evaluation in terms of the number of blocks in serial (a), the number of data flows in parallel (b), and the number of inputs/outputs (c)

overhead in terms of the number of blocks in a data flow (8a) and the number of data flows in parallel (8b), covering many possible variations of the general structure derived in figure 7. We also analyzed how the number of inputs/outputs of a block affects the overhead (figure 8c).

These structures are composed by three kinds of blocks. The *Timestamp source* block generates samples which consist of timestamps that represent the time when the sample was generated. The *dummy blocks* are empty blocks that just propagate their inputs to their outputs. After being generated, the samples will go through the *dummy block* chain. When a sample arrives at the *Timestamp sink* block, the timestamp is compared with the current time, obtaining the time the sample took to go through the *dummy block* chain. Since the *dummy blocks* are empty, this resulting time represents only the overhead imposed by the architecture on the data flow. There is also the *Fork block* and *Join block* which are used to fork and join the data flows, respectively.

Table 1: Configuration parameters used on the experiments

Parameter	Value
Synthesis tool	ISE/EDK 10.1
Compiler	GCC 4.0.2
FPGA clock	100 MHz
Microprocessor clock	100 MHz
α	1
Num. of input ports	32
Num. of output ports	32
HW FIFO size	8 bits wide with 16 elements
Num. of FIFOs	64

Table 2: Amount of HW resources used by the synthesized structures

Resource	Our IPs	EDK IPs	Full System
4-input LUTs	37%	35%	72%
Slice Flip Flops	67%	31%	98%
Occupied Slices	70%	55%	99%
RAM blocks	0%	63%	63%
Max. frequency	167 MHz	109 MHz	107 MHz

5.2 System Implementation and Configuration

To evaluate these three structures, we have implemented HYRA on the Xilinx’s ML403 Embedded Platform [29]. The ML403 features a Virtex-4 FPGA with an embedded PowerPC 405 microprocessor. In order to use the same hardware configuration, we have synthesized the hardware with all of the necessary blocks for all experiments. Table 1 shows the parameters used in the experiment’s setup. The α factor was fixed to 1 in order to provide an evaluation considering low jitter requirements. The last four parameters are related to the configuration of the interconnect structure in hardware.

Table 2 shows the resource consumption of the generated hardware. Separate results are shown for HYRA’s structures along with the HW dummy blocks, and for the system IPs generated by EDK (internal memory, memory controller, interruption controller, UART, etc). Due to the lack of memory blocks available on the device, we chose to implement the FIFOs using the SRL16 capabilities to convert a 4-input LUT into a 16-bit shift register. This feature not only allows us to save memory blocks, but also yields high performance and low cost FIFO implementation [30].

Our architecture alone uses about 65% of the available device resources. This apparently high resource usage is due to the very limited amount of logic available on the device used. When compared to other system IPs, we can see that HYRA uses slightly more resources than a complete set of basic IO and memory peripherals. Also, if we consider that a mid-end Xilinx Spartan-6 [31] FPGA has about 10 times more logic resources than the Virtex-4 device in the ML403 platform, our architecture imposes negligible area overhead in more recent devices.

5.3 Performance Overhead

We have done tests to determine the performance overhead that we defined as the latency between the *Times-tamp source and sink* blocks in the three basic data flow structures. We have implemented each structure in hardware and software and executed tests with the number of dummy blocks ranging from 1 to 16. In each test 6×10^7 samples were generated and we obtained the average value of the latencies of each sample and the standard deviation. We used the standard deviation to obtain the coefficient of variation which yields normalized values for the overhead variation among the various configurations. A sampling rate of 1×10^6 samples/second was used in the tests with blocks in hardware. For the tests with blocks in software we used a sampling rate of 1×10^4 due to the low speed of the PowerPC processor.

Figures 9 and 10 show the results. When using only software blocks, the overhead grows linearly in relation to the increase in the number of blocks and the number of inputs and outputs for all structures. The coefficient of variation remained low in all configurations, and the lowest values for the multiple input/output dummy block configurations were due to the constant number of threads in the system. When using only HW blocks, the latency was about four orders of magnitude lower than when using software blocks and, as expected, except for the serial block configuration, the latency remained constant regardless the size of the structure, due to the full parallelism that can be explored in this kind of architecture. There is also a null coefficient of variation in the hardware operations.

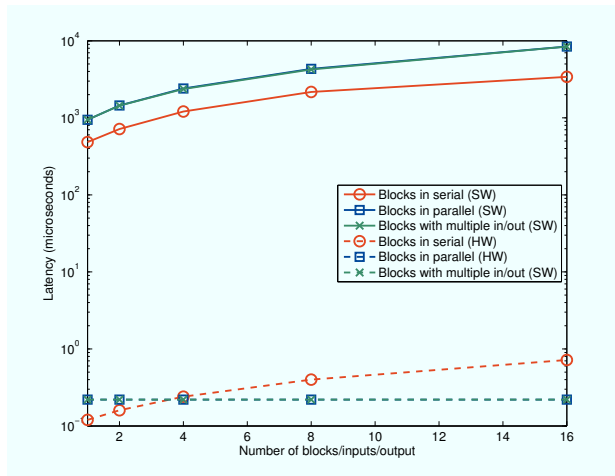


Figure 9: Average latency for blocks in serial, in parallel and with multiple input/output




Figure 10: Coefficient of variation of the proposed architecture latency

To evaluate de latency in the communication between components implemented in HW and components

implemented in SW, we used the data flow shown in figure 11 which cover operations on both *SW-HW FIFOs* and *HW-SW FIFOs*. We evaluated both source/sink components in SW and HW. Figure 12 shows the results of the same tests described previously executed on this two structures, and compares with the results obtained for SW and HW only components. Both interleaved data flows yielded similar results. Data flow (a) have more SW blocks then (b), thus showing a higher SW management overhead. However, both have two *SW-HW FIFOs* and two *HW-SW FIFOs* connecting the blocks, which shows that the read/write operations in the channels represents the most significant overhead.

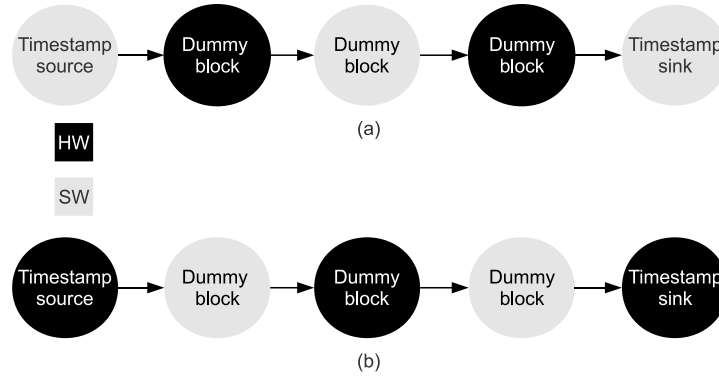


Figure 11: Serial data flows with interleaved SW and HW blocks

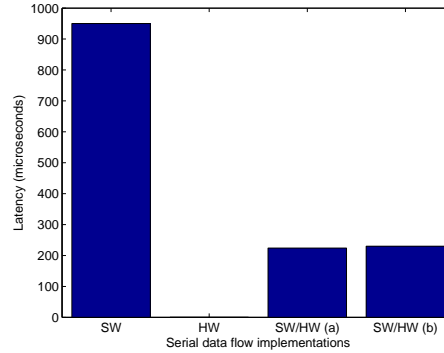


Figure 12: Average latencies on the serial data flow with interleaved SW and HW blocks

We also compared the overhead of our architecture to GNU Radio. For this comparison, we replicated the same tests described previously using GNU Radio running over a Linux operating system in a PC. Our architecture and EPOS were compiled for the IA32 architecture only with software blocks support, and evaluated in the same system. For the GNU Radio experiment, we used GNU Radio 3.2.2 running on a Linux kernel 2.6.28 and, in order to avoid interference from other Linux processes, the GNU Radio

application was executed using the *enable real-time scheduling* feature, which gives the highest priority to the GNU Radio process. The result for the serial blocks data flow structure shown in figure 13 demonstrates that our architecture performance surpasses GNU Radio between 2 and 4 times, and this difference increases as the number of blocks in the processing chain increases. Figure 14 also shows that we are able to achieve smaller latency variations as well.




Figure 13: Average latency of the proposed architecture VS GNU Radio on the serial blocks data flow structure




Figure 14: Coefficient of variation of the proposed architecture latency VS GNU Radio on the serial blocks data flow structure

5.4 Discussion

The results show that our architecture yields superior performance than an equivalent, in terms of abstraction level, commonly used architecture. Only software components were used in this comparison, if a hardware device was used to generate the timestamps, GNU Radio would suffer an additional disadvantage. In GNU Radio, the use of any hardware device to obtain or sink data from/to the environment requires Linux drivers whose performance is mostly limited by the kernel's abstraction layer. A previous work [6] shows that, due to the Linux kernel overhead, the standard deviation of the time a sample takes to get to the processing chain after being generated in the RF Front-end is higher than the average time. This problem does not appear in EPOS mainly for two reasons. First, EPOS is application specific: only the components required by the application are included in the system, and all of OS resources are dedicated towards a single application. Second, the metaprogrammed hardware mediators are dissolved within the application when the system is compiled, which leads to higher performance. Also, the virtually null variation obtained when hardware

blocks are used in a dedicated platform (ML403) is especially important when implementing time-strict protocols, like TDMA based protocols.

However, even with known latency problem, the GNU Radio is widely used and several protocols have been successfully implemented using it. The results have shown that with our architecture we were able to bring similar functionality with superior performance to the embedded system domain, which leads to the conclusion that our architecture is suitable for the implementation of high-end protocols in embedded systems.

6 Conclusion

In this paper we have introduced HYRA, an *Hybrid Radio Architecture* that explores the *Hybrid Component* concept within ADESD to enable the implementation of SDRs as direct mappings of high-level SDF models. Each SDR functional block in the SDF model is implemented as an hybrid component that can be subsequently plugged into HYRA's framework. As hybrid components, HYRA SDR blocks can be implemented as arbitrary combination of software and hardware on FPGA-based platforms. The programmable interconnect infrastructure in HYRA's framework ensures transparency in this respect. FIFO channels can be fine tuned to fulfill the requirements of a given SDR protocol, while the controller dynamically coordinates the flow of data between components.

In comparison with other approaches, HYRA addresses the implementation of SDRs in the context of embedded systems from a higher level of abstraction. Yet the evaluation results presented in this paper confirm that the overhead caused by the proposed architecture in terms of latency and general resource consumption is much smaller than that of GNU Radio, a widely accepted architecture. Furthermore, our experiments demonstrated that HYRA can be implemented on reconfigurable hardware platform with minimal additional resources. In combination, this factors confirm that our architecture meet the requirements for the implementation of high-end protocols in embedded systems.

References

- [1] E. Buracchini, "The software radio concept," *IEEE Communications Magazine*, vol. 38, no. 9, pp. 138–143, 2000.

- [2] E. A. Lee and D. G. Messerschmitt, "Static Scheduling of Synchronous Data Flow Programs for Digital Signal Processing," *IEEE Transactions on Computers*, vol. C-36, pp. 24–35, 1987.
- [3] S. Choi, R. Scrofano, V. K. Prasanna, and J.-W. Jang, "Energy-efficient signal processing using FPGAs," in *FPGA '03: Proceedings of the 2003 ACM/SIGDA eleventh international symposium on Field programmable gate arrays*. New York, NY, USA: ACM, 2003, pp. 225–234.
- [4] A. C. H. Ng, J. W. Weijers, M. Glassee, T. Schuster, B. Bougard, and L. Van der Perre, "ESL design and HW/SW co-verification of high-end software defined radio platforms," in *CODES+ISSS '07: Proceedings of the 5th IEEE/ACM international conference on Hardware/software codesign and system synthesis*. New York, NY, USA: ACM, 2007, pp. 191–196.
- [5] H. Marcondes and A. A. Fröhlich, "A Hybrid Hardware and Software Component Architecture for Embedded System Design," in *Proceedings of the International Embedded System Symposium*, Langenargen, Germany, 2009, pp. 259–270.
- [6] G. Nychis, T. Hottelier, Z. Yang, S. Seshan, and P. Steenkiste, "Enabling MAC Protocols Implementation on Software-defined Radios," in *Networked Systems Design and Implementation*, 2009.
- [7] GNU FSF project, "The GNU Radio," 2010. [Online]. Available: <http://www.gnu.org/software/gnuradio>
- [8] Ettus Research, "Usrp," 2010. [Online]. Available: <http://www.ettus.com>
- [9] J. Glossner, E. Hokenek, and M. Moudgill, "The Sandbridge Sandblaster Communications Processor," in *3rd Workshop on Application Specific Processors*, 2004, pp. 53–58.
- [10] J. Glossner, D. Iancu, M. Moudgill, G. Nacer, S. Jinturkar, and M. Schulte, "The sandbridge SB3011 SDR platform," in *SympoTIC '06: Joint IST Workshop on Mobile Future and the Symposium on Trends in Communications.*, 24-27 2006, pp. ii –v.
- [11] Y. Lin, H. Lee, M. Woh, Y. Harel, S. Mahlke, T. Mudge, C. Chakrabarti, and K. Flautner, "SODA: A Low-power Architecture For Software Radio," in *ISCA '06: 33rd International Symposium on Computer Architecture*, 2006, pp. 89–101.

- [12] M. Woh, Y. Lin, S. Seo, S. Mahlke, T. Mudge, C. Chakrabarti, R. Bruce, D. Kershaw, A. Reid, M. Wilder, and K. Flautner, “From SODA to scotch: The evolution of a wireless baseband processor,” in *MICRO 41: Proceedings of the 41st annual IEEE/ACM International Symposium on Microarchitecture*. Washington, DC, USA: IEEE Computer Society, 2008, pp. 152–163.
- [13] K. van Berkel, F. Heinle, P. P. E. Meuwissen, K. Moerman, and M. Weiss, “Vector processing as an enabler for software-defined radio in handheld devices,” *EURASIP: Journal on Applied Signal Processing*, vol. 2005, pp. 2613–2625, 2005.
- [14] Steven Kelem, Brian Box, Stephen Wasson, Robert Plunkett, Joseph Hassoun, and Chris Phillips, “An Elemental Computing Architecture for SD Radio,” in *SDR ’07: 2007 Software Defined Radio Technical Conference*, 2007.
- [15] Y. Lin, M. Kudlur, S. Mahlke, and T. Mudge, “Hierarchical coarse-grained stream compilation for software defined radio,” in *CASES ’07: Proceedings of the 2007 international conference on Compilers, architecture, and synthesis for embedded systems*. New York, NY, USA: ACM, 2007, pp. 115–124.
- [16] Y. Choi, Y. Lin, N. Chong, S. Mahlke, and T. Mudge, “Stream Compilation for Real-Time Embedded Multicore Systems,” in *CGO ’09: Proceedings of the 7th annual IEEE/ACM International Symposium on Code Generation and Optimization*. Washington, DC, USA: IEEE Computer Society, 2009, pp. 210–220.
- [17] Mentor Graphics, “Catapult C Synthesis,” 2010. [Online]. Available: http://www.mentor.com/products/esl/high_level_synthesis/
- [18] Z. Guo, W. Najjar, and B. Buyukkurt, “Efficient hardware code generation for FPGAs,” *ACM Transactions on Architecture and Code Optimization*, vol. 5, no. 1, pp. 1–26, 2008.
- [19] R. Leupers, “Code generation for embedded processors,” in *ISSS ’00: Proceedings of the 13th international symposium on System synthesis*. Washington, DC, USA: IEEE Computer Society, 2000, pp. 173–178.
- [20] F. Plavec, Z. Vranesic, and S. Brown, “Towards compilation of streaming programs into FPGA hardware,” in *FDL 2008. Forum on Specification, Verification and Design Languages*, 2008, pp. 67–72.

- [21] A. A. Fröhlich, “Application-Oriented Operating Systems,” Ph.D. dissertation, Technical University of Berlin, Berlin, 2001.
- [22] P. P. Pande, C. Grecu, A. Ivanov, and R. Saleh, “Design of a switch for network on chip applications,” in *ISCAS '03. Proceedings of the 2003 International Symposium on Circuits and Systems*, vol. 5, 2003, pp. V217–V220.
- [23] N. Bellas, S. M. Chai, M. Dwyer, and D. Linzmeier, “Mapping streaming architectures on reconfigurable platforms,” *SIGARCH Comput. Archit. News*, vol. 35, no. 3, pp. 2–8, 2007.
- [24] D. Gomez and J. Gonzalez Villarruel, “Analysis of a Software Bluetooth Modem Based on a DSP Implementation,” in *CONIELECOMP 2005: Proceedings of the 15th International Conference on Electronics, Communications and Computers*, 28-02 2005, pp. 108–112.
- [25] C. Ghosh and D. Agrawal, “Cross Layer Performance Evaluation of a Software Defined Radio UWB Receiver using Turbo Decoding,” in *SECON '07. 4th Annual IEEE Communications Society Conference on Sensor, Mesh and Ad Hoc Communications and Networks*, 2007, pp. 639–646.
- [26] T. Schmid, T. Dreier, and M. B. Srivastava, “Software Radio Implementation of Shortrange Wireless Standards for Sensor Networking,” in *Conference On Embedded Networked Sensor Systems*, 2006.
- [27] H. Eslami, G. Patel, C. P. Sukumar, S. V. Tran, A. M. Eltawil, R. Rao, and C. Dick, “Demonstration of highly programmable downlink OFDMA (WiMax) transceivers for SDR systems,” in *MobiHoc '09: Proceedings of the tenth ACM international symposium on Mobile ad hoc networking and computing*. New York, NY, USA: ACM, 2009, pp. 325–326.
- [28] A. Kountouris, C. Moy, L. Rambaud, and P. Le Corre, “A reconfigurable radio case study: a software based multi-standard transceiver for UMTS, GSM, EDGE and Bluetooth,” in *VTC 2001: IEEE VTS 54th Vehicular Technology Conference*, vol. 2, 2001, pp. 1196–1200.
- [29] “Virtex-4 ML403 Embedded Platform,” 2010. [Online]. Available: <http://www.xilinx.com/products/-devkits/HW-V4-ML403-UNI-G.htm>
- [30] K. Chapman, “Saving Costs with the SRL16E,” Xilinx, Tech. Rep., May 2008.

[31] Xilinx, “Spartan-6 FPGA Family.” [Online]. Available: <http://www.xilinx.com/products/spartan6-/index.htm>